

CS 455: INTRODUCTION TO DISTRIBUTED SYSTEMS [HDFS]

Shrideep Pallickara
Computer Science
Colorado State University

March 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.1

Frequently asked questions from the previous class survey

- How is block location reconstructed on startup?
- Difference between Hadoop fs and HDFS??
- Are namenodes typically on higher quality nodes?
- HDFS block report interval: `dfs.blockreport.intervalMsec = 21600000`

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.2

Topics covered in this lecture

- Hadoop Distributed File System
 - Writing Data
 - Replication
 - Data integrity
 - Parallel Copying
 - Coherency Model

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.3

WRITING DATA

March 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.4

File writes


- We will look at creating a new file and writing data to it
- File creation is done using `create()` on `DistributedFileSystem`
- `DistributedFileSystem` does an **RPC** to the namenode
 - Namenode checks existence of file and permissions
 - Creates file in the filesystem's namespace with no blocks in it

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.5

Data flow in HDFS [writes]



March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.6

Anatomy of a file write

- DistributedFileSystem returns an FSDataOutputStream for client to write data to
- FSDataOutputStream wraps a DFSOutputStream
 - DFSOutputStream handles communications with the datanodes and the namenode

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.7

As the client writes data ...

- DFSOutputStream splits it into **packets**
 - Written to an internal queue, the **data queue**
- Data queue is consumed by the DataStreamer
- DataStreamer asks namenode to allocate new blocks
 - Pick list of suitable datanodes to store replicas
 - List of datanodes forms a **pipeline**

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.8

Assuming a replication level of 3

- DataStreamer streams packets to the first datanode in the pipeline
 - 1st datanode stores the packet and forwards it to the 2nd datanode in pipeline
- The second datanode stores the packet and forwards it to the 3rd (and last) datanode in pipeline

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.9

Managing acknowledgements

- DFSOutputStream maintains an internal queue of packets waiting to be ACKed by datanodes
 - This is the **ack queue**
- When is a packet removed from the ACK queue?
 - Only when it has been acknowledged by all the datanodes in the pipeline

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.10

Handling datanode failures during writes [1/2]

- The pipeline is closed
- Any packets in the ack queue are added to the front of the data queue
 - Nodes downstream from the failed node will not miss any packets
- Current block on good datanodes is given a new identity
 - Allows **partial block on failed node** to be **deleted** if that datanode recovers later on

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.11

Handling datanode failures during writes [2/2]

- Failed datanode is removed from the pipeline
- **Remainder** of the block's data is written to the two good datanodes in the pipeline
- Namenode **notifies** block is **under-replicated**
 - Arranges for replicas to be created on another node
- Subsequent blocks are treated as normal

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.12

It is possible that multiple datanodes fail while a block is being written

- As long as `dfs.replication.min` (default 1) replicas are written, the write will succeed
- Block is **asynchronously replicated** across cluster until its target replication factor is reached
 - `dfs.replication` (default 3)

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.13

When a client has finished writing data

- It calls `close()` on the stream
- **Flushes** all remaining packets to the datanode pipeline
 - Wait for acknowledgements before contacting the namenode to signal that file is complete
- Namenode knows about blocks that comprise the file
 - `DataStreamer` requests block allocations
 - Client only waits for blocks to be minimally replicated

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.14

REPLICA PLACEMENTS

March 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.15

Replica placement

[1/2]

- **Trade-off** between reliability, read bandwidth, and write bandwidth
- Placing all replicas on a single node?
 - Lowest write bandwidth penalty since replication pipeline runs on a single node
 - Offers **no redundancy**

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.16

Replica placement

[2/2]

- Read bandwidth is high for off-rack reads
- Placing replicas in different data centers
 - Maximizes redundancy at the cost of bandwidth

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.17

Default replication strategy in Hadoop

- Place **first** replica *on the same node* as the client
 - If client runs outside the cluster, 1st node is chosen at random
- The **second** replica is placed *on a different rack* from the first
 - Chosen at random
- **Third** replica is placed *on the same rack as the second*
 - Different node is chosen at random
- **Further** replicas are placed *on random nodes* in the cluster
 - Avoid placing too many replicas on the same rack

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.18

Default strategy balances

- **Reliability**
 - ▢ Blocks are stored on different racks
- **Write bandwidth**
 - ▢ Writes traverse a single network switch
- **Read bandwidth**
 - ▢ Choice of two racks to read from
- **Block distribution** across cluster
 - ▢ Clients write a single block on the local rack

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.19

Once the replica locations have been chosen

- A pipeline is built
- Pipeline takes network topology into account

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.20

COHERENCY MODEL

March 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.21

A quick look at assertThat in JUnit

- **Format**
 - ▢ `assertThat([value], [matcher statement]);`
- **Examples**
 - ▢ `assertThat(x, is(3));`
 - ▢ `assertThat(x, is(not(4)));`
 - ▢ `assertThat(responseString, either(containsString("color")).or(containsString("colour")));`
 - ▢ `assertThat(myList, hasItem("3"));`

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.22

Assertion syntax

- Readable
- Think in terms of **subject, verb, and object**
 - ▢ Assert "x is 3"
- Matcher statements can be negated, combined, or mapped to a collection

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.23

Coherency Model

- For a filesystem, **coherency** describes data **visibility** of reads and writes to a file
- HDFS trades-off some POSIX requirements for performance

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.24

Creation of a file

- After creation, it is visible in the file namespace

```
Path p = new Path("p");  
fs.create(p);  
assertThat(fs.exists(p), is(true));
```

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.25

Contents written to the newly created file

- **Not guaranteed** to be visible
- Even if the stream is flushed
 - File may appear to have length of 0

```
Path p = new Path("p");  
OutputStream out = fs.create(p);  
out.write("content".getBytes("UTF-8"));  
out.flush();  
assertThat(fs.getFileStatus(p).getLen(), is(0L));
```

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.26

Visibility of blocks during writes

- Once more than a block of data is written?
 - The first block is visible
- In general, the current block that is *being written to* is **not visible** to other readers

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.27

The HDFS sync method

- Forces all buffers to be **synchronized** to the datanodes
- After sync () returns successfully?
 - All data written up to that point in the file is persisted and visible to all clients

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.28

When to call sync ()

- With no calls to sync ()
 - Possible to lose up to a block of data due to client or system failure
- However invocations of sync () do have **overheads**
 - **Trade-off** between data robustness and throughput
- Frequency of sync () is application dependent

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.29

PARALLEL COPYING

March 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.30

Parallel copying with distcp

- Enables copying large amounts of data to and from the Hadoop filesystem in **parallel**

```
% hadoop distcp hdfs://namenode1/foo hdfs://namenode2/bar
```

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.31

distcp is implemented as a MapReduce job

- Copying is done by Maps that run in **parallel** across the cluster
 - There are no reducers
- Deciding the number of maps
 - Give each map sufficient data to **minimize overheads** during **task setup**
 - This is specified using the **-m** argument to distcp

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.32

Keeping an HDFS cluster balanced

- HDFS works best when file blocks are **evenly spread** across the cluster
- We need to ensure that distcp does not disrupt this feature
- If we are transferring 1000 GB?
 - Specifying **-m 1** would mean that a single map would do the copy
 - Will be slow
 - The first replica of each block would reside on the node running map (till the disk fills up)

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.33

DATA INTEGRITY

March 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.34

Data Integrity

- I/O operations on disk or network carry a small chance of introducing errors
- With voluminous data movements the chances of data corruption become high
- Checksums
 - Data is **corrupt** if there is a **mismatch** between the original and the newly computed checksum
 - There is also a **small chance** that the checksum is corrupt

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.35

Data integrity in HDFS

- Datanodes are responsible for **verifying** received data before storing the data and checksum
- When clients read data from the datanode, they verify the checksum
 - Compare with checksum stored at the datanode

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.36

DataBlockScanner

- Each datanode runs a DataBlockScanner in the background **periodically**
- **Verifies all blocks** stored on the datanode
- Guards against corruption due to **bit rot** in the physical storage media

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.37

Dealing with corrupted data blocks

- **Heal** corrupted blocks
 - By copying one of the good replicas to produce a new, uncorrupt replica
- When a client detects an error while reading block?
 - Report *both* the bad block and datanode it was reading from
 - Throw `ChecksumException`

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.38

Dealing with corrupted data blocks

- Namenode marks the block replica as **corrupt**
 - Does not direct clients to it
 - Does not try to copy replica to another datanode
- **Schedules a copy** of the block to be replicated on another datanode
 - Restore replication level for the block
- Corrupt replica is then **deleted**

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.39

Disabling checksum

- Useful if you have a corrupt file that you would like to **inspect**
- Pass `false` to `verifyChecksum()` on `FileSystem` before using `open()` to read the file
- From the shell, use the `-ignoreCrc` option with the `-get` or the `-copyToLocal` command

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.40

Client side checksumming

- Done by the Hadoop `LocalFileSystem`
- When you write a file `filename`
 - The filesystem client creates a hidden file `.filename.crc` in the same directory
 - Contains checksums for each chunk of the file
 - Chunk size is stored in the `.crc` file
- Disable checksums when underlying filesystem supports this natively
 - Use `RawLocalFileSystem` instead of `LocalFileSystem`

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.41

COMPRESSION

March 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.42

Compression

- Reduces **space** needed to store files
- Speeds up data **transfers**
 - Across network
 - Disk I/O

March 20, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L17.43

Compression formats that can be used with Hadoop

Compression format	Tool	Algorithm	Filename extension	Splittable?
DEFLATE	N/A	DEFLATE	.deflate	No
Gzip	Gzip	DEFLATE	.gz	No
Bzip2	Bzip2	Bzip2	.bz2	Yes
LZO	Lzop	LZO	.lzo	No*
Snappy	N/A	Snappy	.snappy	No



March 20, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L17.44

Compression Algorithms

- Exhibit a **space-time** trade-off
 - Faster compression/decompression speeds usually result in smaller space savings
- Tools give some control over this trade-off at compression time
 - 9 different options
 - -1 means optimize for speed
 - -9 means optimize for space

March 20, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L17.45

Compression characteristics

- gzip is a **general purpose** compressor
 - Middle of the space/time trade-off
- bzip2 compresses more effectively than gzip
 - But it is slower
 - bzip2 decompression speed is faster than its compression speed
 - But slower than other formats still
- LZO and Snappy optimize for speed
 - Order of magnitude faster but less effective compression than gzip

March 20, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L17.46

A **codec** is the implementation of a compression-decompression algorithm in Hadoop

Compression format	Hadoop CompressionCodec
DEFLATE	org.apache.hadoop.io.compress.DefaultCodec
gzip	org.apache.hadoop.io.compress.GzipCodec
bzip2	org.apache.hadoop.io.compress.Bzip2Codec
LZO	com.hadoop.compression.lzo.LzopCodec
Snappy	org.apache.hadoop.io.compress.SnappyCodec

March 20, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L17.47

CompressionCodec

- To compress data being written to an output stream
 - Use codec.createOutputStream(OutputStream out)
- To decompress data being read from an input stream
 - Use codec.createInputStream(InputStream in)

March 20, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L17.48

Using compression

```
public class StreamCompressor {  
    public static void main(String[] args) throws Exception {  
        String codecClassname = args[0];  
        Class<?> codecClass = Class.forName(codecClassname);  
        Configuration conf = new Configuration();  
        CompressionCodec codec = (CompressionCodec)  
            ReflectionUtils.newInstance(codecClass, conf);  
        CompressionOutputStream out =  
            codec.createOutputStream(System.out);  
        IOUtils.copyBytes(System.in, out, 4096, false);  
        out.finish();  
    }  
}
```

Compresses data read from standard input and writes it to standard output

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.49

Compression and input splits

- Let's look at an uncompressed file stored in HDFS
 - With an HDFS block size of 64 MB, a 1 GB file is stored as 16 blocks
 - MapReduce job will create 16 input splits
 - Processed **independently** as separate map tasks

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.50

If the gzip compressed file is 1 GB

- HDFS stores files as 16 blocks
- Creating a split for each block does not work
 - Impossible to start reading at an **arbitrary block** in the zip stream
 - Impossible for map task to read its split **independently of others**

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.51

Storing gzipped streams

- Gzip uses DEFLATE, which stores data as a **series of compressed blocks**
- The **start of each block is not distinguished** in a way that allows:
 - Reader positioned at arbitrary point in stream to advance to the beginning of the next block
 - There is **no self-synchronizing** with the stream
 - Gzip does not support splitting

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.52

HDFS does not split gzip files

- Single map will process 16 HDFS blocks
- Most of these blocks will not be local to the map
 - Loss of locality
 - Job is not granular ... takes much longer to run

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.53

The same story plays out if you were dealing with LZO files, but ...

- It is possible to **preprocess** LZO files using an indexer tool
- Build an **index** of split points

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.54

Bzip2

- This does provide a **synchronization marker** between blocks
 - 48-bit approximation of pi
- The marker is used to support splitting

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.55

Dealing with large, unbounded files [Log files]

- ① Store the files uncompressed
- ② Use compression format that supports
 - Splitting: Bzip2
 - Indexing to support splitting: LZO
- ③ Split the file into chunks in the application and compress each chunk separately
 - Choose chunk sizes such that the **compressed chunks** are approximately the size of an HDFS block

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.56

Using compression in MapReduce

- To compress the output of MapReduce job
 - In the job config set `mapred.output.compress` property to true
 - Use `mapred.output.compression.codec` to specify the codec
- Alternatively, we can do this using the `FileOutputFormat`

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.57

Using the FileOutputFormat

```
public class MaxTemperatureWithCompression {  
    public static void main(String[] args) throws Exception {  
        Job job = Job.getInstance();  
        job.setJarByClass(MaxTemperature.class);  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        FileOutputFormat.setCompressOutput(job, true);  
        FileOutputFormat.setOutputCompressorClass(job, GzipCodec.class);  
  
        job.setMapperClass(MaxTemperatureMapper.class);  
        job.setCombinerClass(MaxTemperatureReducer.class);  
        job.setReducerClass(MaxTemperatureReducer.class);  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.58

Main reason why Hadoop does not use Java Serialization

- Deserialization creates new instance of each object being deserialized
- Writable objects can be (and are often) reused
- Large MapReduce jobs often serialize/deserialize billions of records
 - Savings from not having to allocate new objects is significant

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.59

The contents of this slide set are based on the following references

- Tom White, *Hadoop: The Definitive Guide*, 3rd Edition, O'Reilly Press. ISBN: 978-1-449-31152-0. Chapters [3 and 4].
- JUnit release notes for version 4.4 available at <http://junit.sourceforge.net/doc/ReleaseNotes4.4.html>

March 20, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L17.60