

CS 455: INTRODUCTION TO DISTRIBUTED SYSTEMS [SPARK]

Shrideep Pallickara
Computer Science
Colorado State University

March 27, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.1

Frequently asked questions from the previous class survey

- 48-bit bookending in Bzip2: does the number have to be "special"?
- Spark seems to have "too many" features/extension libraries??
 - Good or bad
- Code inspection?

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.2

Topics covered in this lecture

- Resilient Distributed Datasets
- Common Transformations and Actions

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.3

RESILIENT DISTRIBUTED DATASET [RDD]

March 27, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.4

Lazy loading allows Spark to see the whole chain of transformations

- Allows it to **compute just the data needed** for the result
- Example:

```
lines = sc.textFile("README.md")
pythonLines = lines.filter(lambda line: "Python" in line)
```
- If Spark were to load and store all lines in the file, as soon as we wrote `lines=sc.textFile()`?
 - Would waste a lot of storage space, since we immediately filter out a lot of lines

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.5

RDD and actions

- RDDs are **recomputed** (by default) every time you run an action on them
- If you wanted to **reuse** an RDD?
 - Ask Spark to **persist** it using `RDD.persist()`
 - After computing it the first time, Spark will store RDD contents in memory (*partitioned* across cluster machines)
 - Persisted RDD is used in future actions

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.6

RDDs: memory residency and immutability implications

- Spark can keep an RDD loaded in-memory on the executor nodes throughout the life of a Spark application for faster access in **repeated computations**
- RDDs are immutable, so **transforming an RDD returns a new RDD** rather than the existing one
- Cross-cutting implications?
 - ▢ Lazy evaluation, in-memory storage, and immutability allows Spark to be easy-to-use, fault-tolerant, scalable, and efficient

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.7

Every Spark program and shell works as follows

- ① **Create** some input RDD from external data
- ② **Transform** them to define new RDDs using transformations like `filter()`
- ③ Ask Spark to **`persist()`** any intermediate RDDs that needs to be reused
- ④ **Launch actions** such as `count()`, etc. to kickoff a parallel computation
 - ▢ Computing is optimized and executed by Spark

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.8

A CLOSER LOOK AT RDD OPERATIONS

March 27, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.9

RDDs support two types of operations

- Transformations
 - ▢ Operations that **return a new RDD**. E.g.: `filter()`
- Actions
 - ▢ Operations that **return a result** to the driver program or write to storage
 - ▢ Kicks off a computation. E.g.: `count()`
- Distinguishing aspect?
 - ▢ Transformations return RDDs
 - ▢ Actions return *some other* data type

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.10

Transformations

- Many transformations are **element-wise**
 - ▢ Work on only one element at a time
 - Some transformations are not element-wise
 - ▢ E.g.: We have a logfile, `log.txt`, with several messages, but we only want to select error messages
- ```
inputRDD = sc.textFile("log.txt")
errorsRDD = inputRDD.filter(lambda x: "error" in x)
```

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.11

## In our previous example ...

- `filter` **does not mutate** `inputRDD`
    - ▢ Returns a pointer to an entirely new RDD
    - ▢ `inputRDD` can still be reused later in the program
  - We could use `inputRDD` to search for lines with the word "warning"
    - ▢ While we are at it, we will use another transformation, `union()`, to print number of lines that contained either
- ```
errorsRDD = inputRDD.filter(lambda x: "error" in x)
warningsRDD = inputRDD.filter(lambda x: "warning" in x)
badlinesRDD = errorsRDD.union(warningsRDD)
```

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.12

In our previous example

- Note how `union()` is different from `filter()`
 - ▢ Operates on 2 RDDs instead of one
- Transformations can actually operate on **any number** of RDDs

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.13

RDD Lineage graphs


- As new RDDs are derived from each other using transformations, Spark *tracks dependencies*
 - ▢ **Lineage graph**
- Uses lineage graph to
 - ▢ Compute each RDD on demand
 - ▢ Recover lost data if part of persistent RDD is lost

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.14

RDD lineage graph for our example



March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.15

Actions

- We can create RDDs from each other using transformations
- At some point, we need to actually **do something** with the dataset
 - ▢ **Actions**
- Forces *evaluations of the transformations* required for the RDD they were called on

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.16

Let's try to print information about `badLinesRDD`

```
print "Input had " + badLinesRDD.count() + " concerning lines"
print "here are 10 examples:"
for line in badLinesRDD.take(10)
    print line
```

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.17

RDDs also have a `collect` to retrieve the entire RDD

- Useful if program filters RDD to a very small size and you want to deal locally
 - ▢ Your entire dataset must fit in memory on a single machine to use `collect()` on it
 - Should NOT be used on large datasets
- In most cases, RDDs **cannot be collect()** ed to the driver
 - ▢ Common to write data out to a distributed storage system ... HDFS or S3

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.18

Lazy Evaluation

- Transformations on RDDs are **lazily evaluated**
 - Spark will not begin to execute until it sees an action
- Uses this to **reduce the number of passes** it has to take over data by grouping operations together
- What does this mean?
 - When you call a transformation on an RDD (for e.g. map) the operation is not immediately performed
 - Spark internally records metadata that operation is requested

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.19

How you should think of RDDs

- Rather than thinking of it as containing specific data
 - Best to think of it as **containing instructions on how to compute the data** that we build through transformations
- Loading data into a RDD is lazily evaluated just as transformations are

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.20

COMMON TRANSFORMATIONS AND ACTIONS

March 27, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.21

Element-wise transformations: filter()

- Takes in a function and returns an RDD that only has elements that pass the filter() function


March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.22

Element-wise transformations: map()

- Takes in a function and applies it to each element in the RDD
- Result of the function is the new value of each element in the resulting RDD



March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.23

Things that can be done with map()

- Fetch website associated with each URL in collection to just squaring numbers
- map()'s return type does not have to be the same as its input type
- Multiple output elements for each input element?
 - Use flatMap()


```
lines=sc.parallelize(["hello world", "hi"])
words=lines.flatMap(lambda line: line.split(" "))
words.first() # returns hello
```

March 27, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L19.24

Difference between map and flatMap



March 27, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L19.25

Pseudo set operations


- RDDs support many of the operations of mathematical sets such as union, intersection, etc.
- Even when the RDDs themselves are not properly sets

March 27, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L19.26

Some simple set operations




March 27, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L19.27

Cartesian product between two RDDs



March 27, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L19.28

COMMON ACTIONS

March 27, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L19.29

Actions on Basic RDDs

- `reduce()`
 - Takes a function that operates on two elements in the RDD; returns an element of the same type
 - E.g. of such an operation? $+$ sums the RDD
- ```
sum = rdd.reduce((x,y) => x + y)
```
- `fold()` takes a function with the same signature as `reduce()`, but also takes a "zero value" for initial call
    - "Zero value" is the **identity element** for initial call
    - E.g., 0 for  $+$ , 1 for  $*$ , empty list for concatenation

March 27, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L19.30

Both `fold()` and `reduce()` require return type of same type as the RDD elements

- The `aggregate()` removes that constraint
  - For e.g. when computing a running average, maintain both the count so far and the number of elements

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.31

## EXAMPLES: BASIC ACTIONS ON RDDs

March 27, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.32

### Examples: Basic actions on RDDs [1/7]

- Our RDD contains {1, 2, 3, 3}
- **`collect()`**
  - Return all elements from the RDD
  - Invocation: `rdd.collect()`
  - Result: {1, 2, 3, 3}

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.33

### Examples: Basic actions on RDDs [2/7]

- Our RDD contains {1, 2, 3, 3}
- **`count()`**
  - Number of elements in the RDD
  - Invocation: `rdd.count()`
  - Result: 4

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.34

### Examples: Basic actions on RDDs [3/7]

- Our RDD contains {1, 2, 3, 3}
- **`countByValue()`**
  - Number of times each element occurs in the RDD
  - Invocation: `rdd.countByValue()`
  - Result: {(1,1), (2,1), (3,2)}

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.35

### Examples: Basic actions on RDDs [4/7]

- Our RDD contains {1, 2, 3, 3}
- **`take(num)`**
  - Return num elements from the RDD
  - Invocation: `rdd.take(2)`
  - Result: {1, 2}

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.36

#### Examples: Basic actions on RDDs [5/7]

- Our RDD contains {1, 2, 3, 3}
- reduce(func)**
  - Combine the elements of the RDD together in parallel
  - Invocation: `rdd.reduce( (x,y) => x + y )`
  - Result: 9

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.37

#### Examples: Basic actions on RDDs [6/7]

- Our RDD contains {1, 2, 3, 3}
- aggregate(zeroValue)(seqOp, combOp)**
  - Similar to `reduce()` but used to return a different type
  - Invocation:
    - `rdd.aggregate( (0,0),`  
`( (x,y) => (x._1 + y, x._2 + 1),`  
`(x,y) => (x._1 + y._1, x._2 + y._2) )`
  - Result: (9, 4)

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.38

#### Examples: Basic actions on RDDs [7/7]

- Our RDD contains {1, 2, 3, 3}
- foreach(func)**
  - Apply the provided function to each element of the RDD
  - Invocation: `rdd.foreach(func)`
  - Result: Nothing

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.39

## PERSISTENCE (CACHING)

March 27, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.40

## Why persistence?

- Spark RDDs are lazily evaluated, and we may sometimes wish to use the same RDD multiple times
  - Naively, Spark will **recompute RDD and all of its dependencies** each time we call an action on the RDD
    - Super expensive for iterative algorithms
- To avoid recomputing RDD multiple times?
  - Ask Spark to **persist** the data
  - The nodes that compute the RDD, store the partitions
  - E.g.: `result.persist(StorageLevel.DISK_ONLY)`

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.41

## Coping with failures

- If a node that has data persisted on it fails?
  - Spark recomputes lost partitions of data when needed
- Also, replicate data on multiple nodes
  - To handle node failures without slowdowns

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.42

## Persistence Levels for Spark

| Level               | Space Used | CPU time | In Memory | On disk | Comments                                                                                              |
|---------------------|------------|----------|-----------|---------|-------------------------------------------------------------------------------------------------------|
| MEMORY_ONLY         | High       | Low      | Y         | N       |                                                                                                       |
| MEMORY_ONLY_SER     | Low        | High     | Y         | N       |                                                                                                       |
| MEMORY_AND_DISK     | High       | Medium   | Some      | Some    | Spills to disk if there is too much data to fit in memory                                             |
| MEMORY_AND_DISK_SER | Low        | High     | Some      | Some    | Spills to disk if there is too much data to fit in memory. Stores serialized representation in memory |
| DISK_ONLY           | Low        | High     | N         | Y       |                                                                                                       |

March 27, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L19.43

## What if you attempt to cache too much data that does not fit in memory?

- Spark will **evict old partitions** using a Least Recently Used Cache policy
  - For memory only storage partitions, it will be recomputed the next time they are accessed
  - For memory\_and\_disk ones? Write them out to disk
- RDDs also come with a method, `unpersist()`
  - Manually remove data elements from the cache

March 27, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L19.44

## WORKING WITH KEY/VALUE PAIRS

March 27, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L19.45

## RDDs of key/value pairs

- Key/value RDDs are commonly used to perform aggregations
  - Might have to do ETL (Extract, Transform, and Load) to get data into key/value formats
- Advanced feature to control layout of pair RDDs across nodes
  - Partitioning**

March 27, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L19.46

## RDDs containing key/value pairs

- Are called **pair RDDs**
- Useful **building block** in many programs
  - Expose operations that allow actions on each key in parallel or regroup data across network
  - `reduceByKey()` to aggregate data separately for each key
  - `join()` to merge two RDDs together by grouping elements of the same key

March 27, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L19.47

## PAIR RDDs

March 27, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L19.48



## Pair RDDs

- RDDs that contain **key/value pairs**
- Expose partitions that allow you to act on each key in parallel or regroup data across the network

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.49

## Creating Pair RDDs

- `pairs=lines.map(lambda x: (x.split(" ")[0], x))`
  - Creates a pairRDD using the first word as the key
- Java does not have a built-in tuple type
  - `scala.Tuple2` class
    - `new Tuple2(elem1, elem2)`

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.50

## The contents of this slide-set are based on the following references

- *Learning Spark: Lightning-Fast Big Data Analysis*. 1st Edition. Holden Karau, Andy Konwinski, Patrick Wendell, and Matei Zaharia. O'Reilly. 2015. ISBN-13: 978-1449358624. [Chapters 1-4]
- Karau, Holden; Warren, Rachel. *High Performance Spark: Best Practices for Scaling and Optimizing Apache Spark*. O'Reilly Media. 2017. ISBN-13: 978-1491943205. [Chapter 2]

March 27, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L19.51