

CS 455: INTRODUCTION TO DISTRIBUTED SYSTEMS [DISTRIBUTED COORDINATION/MUTUAL EXCLUSION]

Shrideep Pallickara
Computer Science
Colorado State University

April 5, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.1

Frequently asked questions from the previous class survey

- Is the streaming data received and buffered on multiple nodes or just one?
- Does 1 batch = 1 RDD?
- Can you (would it be worth it to) call `persist()` on an RDD in Spark Streaming?

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.2

Topics covered in this lecture

- Distributed Coordination
- Distributed Mutual Exclusion

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

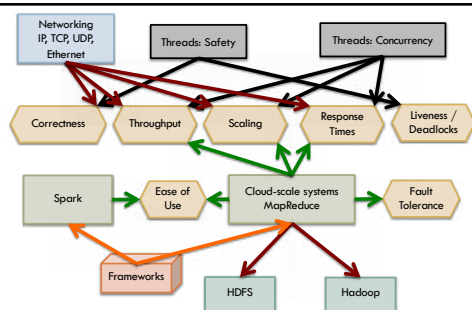
L22.3

THE JOURNEY SO FAR

April 5, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.4



5

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

DISTRIBUTED COORDINATION

April 5, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.6

What we will cover

- Collection of algorithms whose goals vary, but share an aim that is fundamental in distributed systems
 - For a set of processes to:
 - **Coordinate** their actions
 - **Agree** on one or more values

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.7

Communication styles

- Asynchronous communications
 - No timing assumptions
 - Synchronous communications have bounds on
 - Maximum message transmission delay
 - Time to execute each step of a process
 - Clock drift rates
- Allows us to use timeouts to detect process crashes

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.8

Coordination & Agreement

- A set of processes need to **coordinate** actions or **agree** on a set of values
- Must be able to do so even when **hierarchical** relationships do not exist
 - E.g.: Controller-Worker where a single point of failure exists

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.9

Example: Spaceship

- Multiple computers
- Computers that control spaceship must agree on several conditions
 - E.g., Status: Proceed or abort mission
- Coordinate access to shared resources
 - Sensors, actuators, etc.

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.10

DISTRIBUTED MUTUAL EXCLUSION

April 5, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.11

Distributed processes often need to coordinate their activities

- If a collection of processes share a set of resources **mutual exclusion** is needed to:
 - Prevent interference
 - Ensure consistency
- This is the critical section problem in OS.

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.12

Distributed mutual exclusion

- Extension to distributed systems of the familiar problem of avoiding race conditions
 - In kernels and multi-threaded applications
- Shared variables or facilities provided by a local kernel cannot be used to solve this
- Solution must be based *solely* on **message passing**

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.13

Distributed mutual exclusion

- Consider a set of N processes p_i $i=1, 2, \dots, N$
 - These **do not share variables**
- Processes access common resources
 - They do so in a critical section

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.14

SUMMARY OF APPROACHES

April 5, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.15

Approaches to distributed mutual exclusion

- Token-based solutions
- Permission-based solutions

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.16

Token-based solutions

- Mutual exclusion is achieved by *passing* a special message (**token**) between the processes
- There is only one token
 - Whoever has that token is allowed to access shared resource
- When finished, token is passed to another process

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.17

Token-based solutions: Advantages

- Depending on how processes are organized, fairly easy to avoid starvation
- Deadlocks can also be avoided

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.18

Token-based solutions: Disadvantages

- When the token is **lost** – for e.g., process holding the token crashes, complex actions need to be taken
- Intricate distributed process needs to be initiated
 - Ensure that a **new** token is created
 - But above all, make sure that that is the **only** token

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.19

Permission-based solutions

- Process wanting to access resource first requires **permission** of **other** processes
- Many different ways to granting this permission

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.20

Structural considerations for the solution

- With a central server
- Without a central server
 - Peer processes must coordinate their accesses to shared resources
 - Occurs routinely on Ethernets and IEEE 802.11 wireless
 - Network interfaces cooperate as peers so that only one node transmits at a time on the shared medium
 - Ethernet: Method of operation "Carrier Sensing, Multiple Access with Collision Detection" or CSMA/CD
 - Wireless: "Carrier Sensing, Multiple Access with Collision Avoidance" CSMA/CA

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.21

ASSUMPTION & REQUIREMENTS

April 5, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.22

Assumptions in our algorithms

- The system is asynchronous
- Processes do not fail
- Message delivery is reliable
 - Delivered eventually and exactly-once

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.23

Application level protocol for entering the critical section

- `enter()`
 - **Block** if necessary
- `resourceAccesses()`
 - **Access shared resources** in the critical section
- `exit()`
 - Allow other processes to enter

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.24

Requirements for distributed mutual exclusion

- ME1: *At most one* process may execute in the critical section at a time
 - **Safety**
- ME2: Requests to enter and exit the critical section *eventually succeed*
 - **Liveness**: Freedom from **deadlocks** and **starvation**
- ME3: If one request *happened-before* another, then entry to the CS is granted in that order

April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.25

Evaluation of the algorithms

- **Bandwidth consumed**
 - Proportional to number of messages sent in each entry and exit operation
- **Client delay** incurred by process for each entry or exit operation
- Effect on **throughput** of the system
 - **Synchronization delay** between one process exiting critical section and next process entering it
 - Throughput is greater when synchronization delay is shorter

April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.26

THE CENTRAL SERVER ALGORITHM

April 5, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.27

The central server algorithm

- Simplest way to achieve mutual exclusion
- Central server *grants authorization* to enter the critical section
- To enter a critical section, process sends request message to the server
 - Awaits reply from server
 - Reply constitutes **token** signifying authorization to enter critical section

April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.28

Acquisition of token


- If no process holds the token?
 - Server replies immediately granting token
- If the token is held by another process?
 - Server does not reply, but *queues the request*
 - When that process exits the critical section, it sends a message giving server back the token
 - If the queue of waiting processes is non-empty, **server chooses oldest entry** in the queue and sends it the token

April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.29

Server managing a mutual exclusion token



April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.30

Evaluating the central server algorithm [1/2]

- Entering critical section
 - Requires 2 messages: **Request** followed by **grant**
 - Delay at the requesting process?
 - Round trip delay
 - There is also the **queuing delay** for messages residing in the queue
- Exiting the critical section requires one **release** message
 - Assuming asynchronous communications means that this does not delay the exiting process

April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.31

Evaluating the central server algorithm [2/2]

- Synchronization delay
 - Release message to server followed by grant to another process: Round trip time
- Server is a **performance bottleneck** for the system
 - Single point of failure as well

April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.32

RING BASED ALGORITHM

April 5, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.33

Ring-based algorithm


- Arrange mutual exclusion between N processes **without** requiring an additional process
- Each process p_i has a communication channel to the next process in the ring, $p_{(i+1) \bmod N}$
- Exclusion is conferred by obtaining a token that is **passed from process to process** in a single direction around the ring
 - E.g. clockwise

April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.34

Ring topology is unrelated to physical connections between underlying nodes



April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.35

Acquisition of token

- When a process that does not need to enter critical section receives the token?
 - Immediately forwards token to its neighbor
- Process that requires token, **waits** until it receives it **and then retains it**
- To **exit** the critical section, process **sends token** to neighbor

April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.36

Properties satisfied by the ring algorithm

- Satisfies ME1 and ME2
- Token is not necessarily acquired in a happened-before manner (ME3)

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.37

Performance analysis

[1/2]

- **Continuously** consumes network bandwidth (except when process is in critical section)
 - Processes send messages around ring even when no process requires critical section entry
- Delay experienced by process requesting entry to critical section?
 - 0: when it has just received the token
 - N messages when it has just passed on the token

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.38

Performance analysis

[2/2]

- Exit from critical section
 - Requires only 1 message
- Synchronization delay between one process' exit and another process' entry into critical section
 - Anywhere between 1 and N message transmissions

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.39

MUTUAL EXCLUSION USING MULTICAST AND LOGICAL CLOCKS RICART & AGARWALA'S ALGORITHM}

April 5, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.40

LOGICAL CLOCKS: If two processes do not interact with each other

- Their clocks **need not** be synchronized
- Lack of synchronization is not observable
 - Does not cause problems

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.41

Lamport's logical clocks

- The **happens-before** relation \rightarrow
- a and b are events in the process; and a occurs before b
 - Then $a \rightarrow b$ is true
- a is event of message sent by one process;
 b is event of message being received in another process
 - Then $a \rightarrow b$ is true

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.42

Some more things about the happens-before relation


- If $a \rightarrow b$ and $b \rightarrow c$, then $a \rightarrow c$
 - **Transitive**
- If events x and y occur in processes that do not exchange messages, then ...
 - $x \rightarrow y$ is not true
 - But, neither is $y \rightarrow x$
 - These events are said to be **concurrent**

April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.43

An example of Lamport's algorithm:




April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.44

An example of Lamport's algorithm:



April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.45

Implementing Lamport's clocks


- ① Before executing an event; P_i executes $C_i \leftarrow C_i + 1$
- ② When P_i sends a message m to P_j ; it sets m 's timestamp $ts(m)$ to C_i in previous step
- ③ Upon receipt of message m , P_j adjusts its own local counter $C_j \leftarrow \max\{C_j, ts(m)\}$ do step (1) and deliver message

April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.46

An application of Lamport's clock: User has \$1000 in bank account initially



April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.47

There is a difference when the orders are reversed

- Our objective for now is consistency
- Both copies must be exactly the same
- Situations like this require **totally-ordered multicast**
 - All messages are delivered in the same order to each receiver
 - Lamport's logical clocks allow us to accomplish this in a completely distributed fashion

April 5, 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L22.48

Using Lamport's clock to order messages

- Process puts received messages into local queue
 - Ordered according to the message's timestamp
- Message can be delivered only if it is **acknowledged** by all the other processes
- If a message is at the head of the queue, and acknowledged by all processes
 - It is delivered and processed

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.49

Other types of logical clocks

- Vector clocks
- Matrix clocks

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.50

The contents of this slide set are based on the following references

- Distributed Systems: Concepts and Design. George Coulouris, Jean Dollimore, Tim Kindberg, Gordon Blair. 5th Edition. Addison Wesley. ISBN: 978-0132143011 [Chapter 15]

April 5, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L22.51