

CS 455: INTRODUCTION TO DISTRIBUTED SYSTEMS [NETWORKING]

Shrideep Pallickara
 Computer Science
 Colorado State University

January 25 2018

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.1

Frequently asked questions from the previous class surveys

- False positives in checksums?
- Frames and bit-patterns
- What if the IP address is changed mid-transmission?
- Payload length: bytes or words (4-bytes)
- Process-per-message: Is there a separate one per-message?
- Why work with IP/TCP/UDP when you can write directly to network?
- Why not send maximum length datagram messages every time?
- Purpose of trailer? End of message or Here's what's next
- Why does OSI need 7 layers?
- Why not use TOS in IPv4 for other things?

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.2

Topics covered in today's lecture

- IP routing
 - IPv6
- UDP
- TCP

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.3

Every network type has a Maximum Transmission Unit (MTU)

- Largest IP datagram that it can carry in its frame
- Smaller than the largest packet-size of network
 - IP datagram needs to fit in the payload of **link-layer frame**

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.4

Fragmentation necessary when datagram path includes network with smaller MTU


- All fragments carry same identifier in **Ident** field
 - To enable fragment reassembly
 - Chosen by the source host
- If all fragments do not arrive at receiving host?
 - ① Receiver **gives up** reassembly [reassembly timeout: 15 seconds RFC0791]
 - ② **Discards** fragments that did arrive
- IP **does not attempt** to recover from missing fragments

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.5

IPv4 Packet header



January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.6

Header fields used in IP fragmentation:
 Fragmentation occurs at 8-byte boundaries

Start of header			
Ident = x	0	DF	Offset = 0
Rest of header			
1400 data bytes			

Unfragmented packet

Start of header			
Ident = x	1	Offset = 0	
Rest of header			
512 data bytes			

Fragmented packet

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.7

Header fields used in IP fragmentation:
 Fragmentation occurs at 8-byte boundaries

Start of header			
Ident = x	1	Offset = 64	
Rest of header			
512 data bytes			

Fragmented packet

Start of header			
Ident = x	0	Offset = 128	
Rest of header			
376 data bytes			

Fragmented packet

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.8

IPv6 (AND COMPARING WITH IPV4)

January 25 2018

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.9

IPv6 versus IPv4: Key Differences

- Source and destination addresses are **128-bits** (16 bytes) in IPv6
- IPv6 treats Options as **extension headers**
- To simplify processing of packets in routers IPv6 **did away with fragmentation**
 - Responsibility for packet fragmentation is at the end points
 - IPv6 hosts must perform : (1) path MTU discovery, (2) perform end-to-end fragmentation, OR (3) send packets no larger than the default MTU=**1280**
- As of 2014, IPv4 still carried >99% of worldwide Internet traffic

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.10

IPv6 Packet Header

0	4	8	12	16	19	31
Version		Traffic Class		Flow Label		
Payload Length				Next Header		Hop Limit
SourceAddr [16 bytes]						
DestinationAddr [16 bytes]						

IPv6 Packet Header is fixed at 40 bytes ... So there is no Header Length

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.11

IPv6 Packet Header: Some more details [1/2]

- Version:** 4 bits [0110]
- Traffic Class:** 6+2 bits
 - Differentiated Services for QoS
 - Anything that ends in 2 "1" bits is intended for experimental or local use
- Flow Label** (20 bits)
 - If it is non-zero: Serves as a hint to routers and switches with multiple outbound paths that these **packets should stay on the same path**, so that they will not be reordered
- Payload length (16 bits):** Size of payload *including* extension headers

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.12

IPv6 Packet Header: Some more details [2/2]


- Next Header (8 bits)
 - Specifies the type of the next header
- Hop Limit (8 bits)
 - Replaces the time-to-live field of IPv4
- Destination and Source Addresses (**128-bits** or 16 bytes each)
- Note: The IPv6 packet **header has no checksum**
 - Transport or application layer protocols are assumed to provide sufficient error detection

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.13

Structure of the IPv6 Packet




PDU typically contains an upper layer protocol header and its payload.
 For e.g.: a TCP segment, UDP Datagram, or an ICMPv6 message

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.14

IPv6 Extension Headers: The chain of pointers using the Next Header field



Fragmentation Header: 44

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.15

DATAGRAM FORWARDING

January 25 2018

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.16

Datagram forwarding in IP: Datagrams contains IP address of destination


- Network part uniquely identifies a single physical network
- Hosts/routers that share the network part
 - Connected to **same** physical network
- Every physical network has a router
 - Connected to at least **one other** physical network

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.17

A simple internetwork: Forwarding table for router R2



January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.18

Example forwarding table:
For Router R2

Network Num	Next Hop
1	R3
2	R1

January 25 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L4.19

Error Reporting in IP communications

- IP drops datagrams when the going gets tough
 - But does not fail silently
- IP always configured with a **companion** protocol
 - Internet Control Message Protocol (ICMP)

January 25 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L4.20

ICMP defines a collection of error messages

- When router/host is unable to process datagrams successfully
 - ICMP error message **sent back to source**
- Examples
 - Destination host is unreachable
 - Reassembly process failed
 - TTL reached 0
 - IP header checksum failed

January 25 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L4.21

ICMP also defines some control messages

- Router sends **control messages** back to host
- Example: **ICMP-Redirect** tells that there is a better route to destination
 - Network has two routers R1 and R2 and host uses R1 as default
 - When R1 receives a datagram and it knows R2 is a better choice?
 - Send ICMP-Redirect to host
 - Host then uses R2 for future datagrams to that host

January 25 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L4.22

UDP
SIMPLE DEMULTIPLEXER

January 25 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L4.23

User Datagram Protocol

- Simplest** possible transport protocol
 - Extends host-to-host into process-to-process communications
- No additional functionality to best-effort service provided by underlying network
- Adds **demultiplexing**
 - Allows applications on a host to **share** the service

January 25 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L4.24

UDP identification of processes


- Processes *indirectly* identify each other
 - Abstract locator called **port**
- Source sends a message to a port
 - Destination receives messages from a port
- Process is identified by a **port on a particular host**

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.25

Format of a UDP header



January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.26

A port is just an abstraction


- Typically implemented as a **message queue**
- When message arrives?
 - Protocol appends message to end of the queue
- **UDP**
 - If the queue is full, message is discarded
 - No flow-control mechanism

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.27

UDP message queue: The port abstraction



January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.28

Some work that UDP does besides demultiplexing: Checksumming

- UDP header
 - Message body
 - **Pseudoheader**: From the IP header
 - Protocol number
 - Source IP address
 - Destination IP address
 - UDP length
 - Used twice
- } Verify if message is delivered between the correct endpoints

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.29

RELIABLE BYTE STREAM TCP

January 25 2018

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.30

Components of Reliable delivery

- **Acknowledgements**
 - Confirm receipt of data
- **Timeouts**
 - *Retransmit* if ACK not received within a specified time
- Use of ACKs and timeouts to implement reliable delivery
 - Sometime called ARQ (*automatic repeat request*)

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.31

Simplest ARQ is the stop-and-wait algorithm


- After transmitting one frame
 - Sender **waits** for ACK before transmitting the next frame
- If the ACK does not arrive after a period of time
 - Sender **retransmits** the original frame

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.32

Stop-and-wait (1/2)




January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.33

Stop-and-wait (2/2)



January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.34

Sliding window: Try to fill the network pipe


- DELAY x BANDWIDTH product is 8 KB
- Data frames = 1KB
- Sender could transmit 9th frame
 - When ACK for the 1st frame arrives

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.35

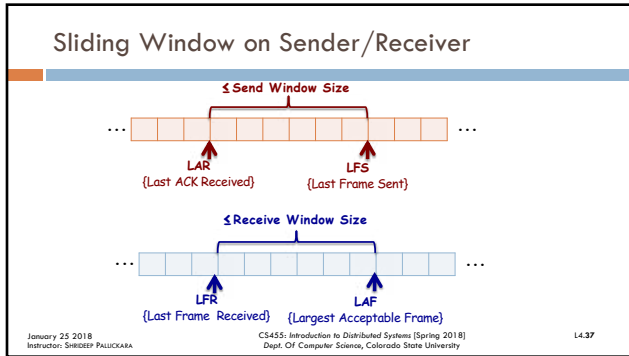
Timeline for the sliding window



January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.36



Transmission Control Protocol (TCP) [1/2]

- **Reliable, in-order** delivery of byte streams
- **Full duplex** protocol
 - Each connection supports a pair of byte streams
 - Flowing in different directions
- Includes **flow control** mechanism
 - Allows receiver to limit the data sender
 - Control how much data can be transmitted at a time

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.38

Transmission Control Protocol (TCP) [2/2]

- Includes **multiplexing** mechanism
 - Multiple apps on a given host
- Implements a **congestion-control** mechanism
 - ① Throttle how fast TCP sends data
 - ② Keep sender from **overloading** the network

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.39

Flow control and congestion control

- **Flow control** is an end-to-end issue
 - Don't overrun capacity of **receiver**
- **Congestion control** is about hosts & networks interact
 - Don't cause **switches** and **links** to be overloaded

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.40

TCP: Setup and Teardown

- Two sides of the connection **agree** to exchange data
 - Establish **shared state**
 - 3 packets exchanged (SYN, SYN-ACK, ACK)
- Connection teardown
 - Let each host know it is OK to **free** the shared state
 - 4 packets exchanged (FIN, ACK, FIN, ACK)

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.41

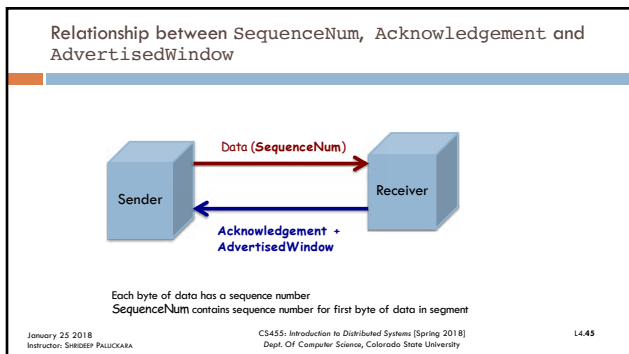
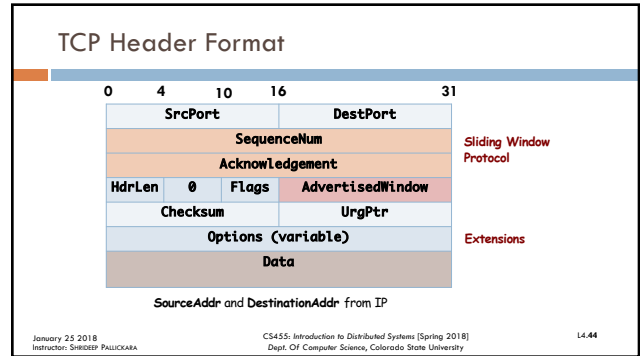
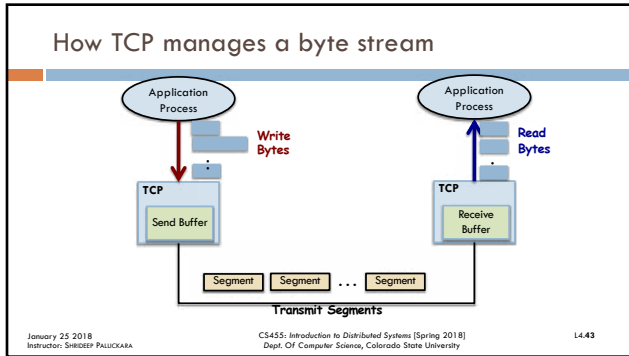
TCP Segments & how they come about

- TCP
 - Accepts data from a data stream
 - Breaks it up into chunks
 - Adds a TCP header ... creating a **TCP segment**
- Segment is then **encapsulated** in a IP datagram
- TCP packet is a term that you will often hear
 - Segment is more precise, packets are generally datagrams, frames are at the link layer

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.42



TCP Sliding Window [1/2]

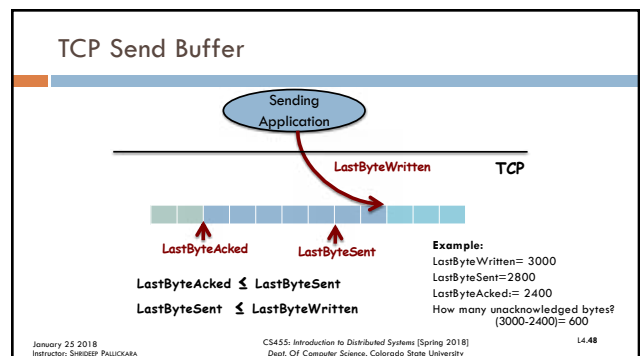
- Guarantees **reliable** delivery of data
- Data is delivered in **order**
- Enforces **flow control** between the sender and receiver

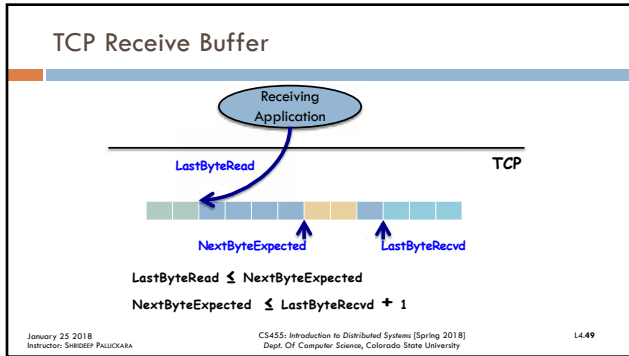
January 25 2018
 Instructor: SHRIDEEP PALICKARA
 CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University
 L4.46

TCP Sliding Window [2/2]

- Sender has a **limit** on unacknowledged data
 - Limited to no more than **AdvertisedWindow** bytes of unacknowledged data
- Receiver **selects** **AdvertisedWindow**
 - Based on memory set aside for connection's buffer space

January 25 2018
 Instructor: SHRIDEEP PALICKARA
 CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University
 L4.47





Flow Control: Buffers are of finite size

MaxSendBuffer and MaxRcvBuffer

- Receiver **throttles** sender
 - Advertises a window
 - No bigger than what it can buffer

$\text{LastByteRcvd} - \text{LastByteRead} \leq \text{MaxRcvBuffer}$

AdvertisedWindow =

$$\text{MaxRcvBuffer} - ((\text{NextByteExpected} - 1) - \text{LastByteRead})$$

Space Utilized in the receiver's buffer

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.50

The advertised window may potentially shrink

- If the process is reading data as fast as it arrives?
 - The advertised window **stays open**
 - i.e. $\text{AdvertisedWindow} = \text{MaxRcvBuffer}$
- If the receiving process falls behind?
 - Advertised window becomes **smaller** with every segment that arrives
 - Until it becomes 0

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.51

Flow Control: Buffers are of finite size

MaxSendBuffer and MaxRcvBuffer

- On the sender size, TCP **adheres** to the advertised window from the receiver

$\text{LastByteSent} - \text{LastByteAcked} \leq \text{AdvertisedWindow}$

EffectiveWindow =

$$\text{AdvertisedWindow} - (\text{LastByteSent} - \text{LastByteAcked})$$

EffectiveWindow should be > 0 before source can send more data

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.52

The contents of this slide-set are based on the following references

- Computer Networks: A Systems Approach. Larry Peterson and Bruce Davie. 4th edition. Morgan Kaufmann. ISBN: 978-0-12-370548-8. Chapters [4, 5]
- <https://en.wikipedia.org/wiki/IPv6>
- Understanding the IPv6 Header:
<https://www.microsoftpressstore.com/articles/article.aspx?p=2225063&seqNum=4>

January 25 2018
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
 Dept. Of Computer Science, Colorado State University

L4.53