

CS 455: INTRODUCTION TO DISTRIBUTED SYSTEMS [THREADS]

Shrideep Pallickara
Computer Science
Colorado State University

February 1, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.1

Frequently asked questions from the previous class survey

- How does the receiver inform the sender that its buffer is no longer full?
- Maximum segment size: Max data that can be received in a single segment
- Retransmission of a lost segment: Do you discard subsequent segments that were received? No.
 - Can the receiver buffer fill up before the "next-in-order" segment arrives?
- How do threads communicate with each other?
- Permanent generation

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.2

Topics covered in this lecture

- Creation and Management
- Thread lifecycle
 - Creating and starting threads
- Stopping and interrupting threads
- Approaches to writing threads
 - Subclassing Threads vs Implementing Runnable

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.3

Thread Abstraction

- A **thread** is a *single execution sequence* that represents a separately schedulable task
 - **Single execution sequence**
 - Each thread executes sequence of instructions – assignments, conditionals, loops, procedures, etc. – just as the sequential programming model
 - **Separately schedulable task**
 - The OS can run, suspend, or resume a thread at any time

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.4

THREADS AND ... HEAPS AND STACKS

February 1, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.5

Threads and heaps

- For performance reasons, heaps may **internally subdivide** their space into per-thread regions
 - Threads can allocate objects at the same time *without interfering* with each other
 - By allocating objects used by the same thread from the same memory region?
 - Cache hit rates may improve
- Each subdivision of the heap has **thread-local variables**
 - Track parts of thread-local heap in use, those that are free, etc.
- New memory allocations (`malloc()` and `new()`) can take memory from **shared heap**, only if local heap is used up

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.6

How big a stack?

- The size of the stack must be large enough to accommodate the **deepest nesting level** needed during the thread's *lifetime*
- Kernel threads
 - ▢ Kernel stacks are allocated in physical memory
 - ▢ The nesting depth for kernels tends to be small
 - ▢ E.g. 8KB default in Linux on an Intel x86
 - ▢ Buffers and data structures are allocated on the heap and never as procedure local variables

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.7

How big a stack?

- User-level stacks are allocated in virtual memory
- To catch program errors
 - ▢ Most OS will trigger **error** if the program stack grows **too large too quickly**
 - Indication of an unbounded recursion
 - ▢ Google's GO will automatically grow the stack as needed ... this is very uncommon
 - ▢ POSIX for e.g. allows default stack size to be library dependent (e.g. larger on a desktop, smaller on a phone)
 - "Exceeding default stack limit is very easy to do, with the usual results"
 - Program termination

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.8

THREAD CREATION

February 1, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.9

Thread creation

- Using the **Thread** class
- Using the **Runnable** interface

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.10

The Thread class

```
package java.lang;

public class Thread implements Runnable {
    public Thread();
    public Thread(Runnable target);
    public Thread(ThreadGroup group, Runnable target);
    public Thread(String name);
    public Thread(ThreadGroup group, String name);
    public Thread(Runnable target, String name);
    public Thread(ThreadGroup group, Runnable target,
                  String name);
    public Thread(ThreadGroup group, Runnable target,
                  String name, long stackSize);

    public void start();
    public void run();
}
```

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.11

Threads require 4 pieces of information

- Thread **name**
 - ▢ Default is Thread-N; N is a unique number
- **Runnable target**
 - ▢ List of *instructions* that the thread executes
 - ▢ Default: run() method of the thread itself
- Thread **group**
 - ▢ A thread is assigned to the thread group of the thread that calls the constructor
- **Stack size**
 - ▢ Store temporary variables during method execution
 - ▢ Platform-dependent: range of legal values, optimal value, etc.

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.12

A simple thread

```
public class RandomGen extends Thread {  
    private Random random;  
    private int nextNumber;  
    public RandomGen() {random = new Random();}  
  
    public void run() {  
        for (;;) {  
            nextNumber = random.nextInt();  
            try {  
  
            } catch (InterruptedException ie) {  
                ... return;  
            }  
        }  
    }  
}
```

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.13

About the code snippet

- Extends the Thread class
- Actual instructions we want to execute is in the run() method
 - Standard method of the Thread class
 - Place where Thread begins execution

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.14

Contrasting the run() and main() methods

- main() method
 - This is where the *first thread starts executing*
 - The *main thread*
- The run() method
 - *Subsequent threads* start executing with this method

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.15

THREAD LIFECYCLE

February 1, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.16

Lifecycle of a thread

- Creation
- Starting
- Terminating
- Pausing, suspending, and resuming

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.17

Thread: Methods that impact the thread's lifecycle

```
public class Thread implements Runnable {  
    public void start();  
    public void run();  
    public void stop();  
    public void resume();  
    public void suspend();  
    public static void sleep(long millis);  
    public boolean isAlive();  
    public void interrupt();  
    public boolean isInterrupted();  
    public static boolean interrupted();  
    public void join();  
}
```

} Deprecated, do not use

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.18

Thread creation

- Threads are represented by instances of the Thread class
- When you extend the Thread class?
 - ▢ Your instances are also Threads
- We looked at the 4 constructor arguments in the Thread class

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.19

Starting a thread

(1/2)

- Thread exists once it's been constructed
 - ▢ But it is *not executing* ... it's in a **waiting** state
- In the waiting state, other threads can *interact* with the existing **thread object**
 - ▢ Object state may be changed by other threads
 - Via method invocations

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.20

Starting a thread

(2/2)

- When we're ready for a thread to begin executing code
 - ▢ Call the **start()** method
 - ▢ start() performs internal house-keeping and *then calls* the **run()** method
- When the start() method returns?
 - ▢ **Two threads** are executing in parallel
 - ① The original thread which just returned from calling start()
 - ② The newly started thread that is executing its run() method

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.21

After a thread's start() method is called

- The new thread is said to be **alive**
- The isAlive() method tells you about the state
 - true: Thread has been started and *is executing* its run() method
 - false: Thread may *not be started* yet or may be *terminated*

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.22

Terminating a thread

- Once started, a thread executes only one method: run()
- This run() may be complicated
 - ▢ May *execute forever*
 - ▢ Call *several other methods*
- Once the run() *finishes* executing, the thread has **completed** its execution

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.23

Like all Java methods, run() finishes when it ...

- ① Executes a return statement
- ② Executes the last statement in its method body
- ③ When it *throws an exception*
 - ▢ Or fails to catch an exception thrown to it

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.24

The only way to terminate a thread?

- Arrange for its `run()` method to **complete**
- But the documentation for the `Thread` class lists a `stop()` method?
 - This has a **race condition** (unsafe), and has been deprecated

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.25

Some more about the `run()` method

- Cannot throw a **checked** exception
- But it can throw an **unchecked** exception
 - Exception that extends the `RuntimeException`
- A thread can be **stopped** by:
 - ① **Throwing** an unchecked exception in `run()`
 - ② **Failing to catch** an unchecked exception thrown by something that `run()` has called

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.26

Pausing, suspending and resuming threads

- Some thread models support the concept of **thread suspension**
 - Thread is told to **pause** execution and then told to **resume** its execution
- Thread contains `suspend()` and `resume()`
 - Suffers from vulnerability to **race conditions**: **deprecated**
- Thread can **suspend its own execution** for a specified period
 - By calling the `sleep()` method

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.27

But sleeping is not the same thing as thread suspension

- With true thread suspension
 - One thread can suspend (and later resume) **another thread**
- `sleep()` affects only the thread that executes it
 - Not possible to tell another thread to go to sleep

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.28

But you can achieve the functionality of suspension and resumption

- Use wait and notify mechanisms
- Threads **must be coded** to use this technique
 - This is **not a generic** suspend/resume that is imposed by another thread

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.29

Thread cleanup

- As long as some other active object holds a reference to the terminated thread object
 - Other threads can execute methods on the terminated thread ... retrieve information
- If the object representing the terminated thread goes **out of scope**?
 - The thread object is **garbage collected**

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.30

Holding onto a thread reference allows us to determine if work was completed

- Done using the `join()` method
- The `join()` method
 - **Blocks** until the thread has completed
 - **Returns immediately** if
 - The thread has already completed its `run()` method
 - You can call `join()` any number of times
- Don't use `join()` to poll if the thread is still running
- Use `isAlive()`

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.31

STOPPING A THREAD

February 1, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.32

Two approaches to stopping a thread

- Setting a flag
- Interrupting a thread

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.33

Stopping a Thread: Setting a flag

- **Set some internal flag** to signal that the thread should stop
- Thread periodically **queries the flag** to determine if it should exit

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.34

Stopping a Thread: Setting a flag

```
public class RandomGen extends Thread {
    private volatile boolean done = false;

    public void run() {
        while (!done) {
            ...
        }
    }

    public void setDone() {
        done = true;
    }
}
```

`run()` method investigates the state of the done variable on every loop.
Returns when the done flag has been set.

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.35

Interrupting a thread

- In the previous slide, there may be a **delay** in the `setDone()` being invoked & thread terminating
 - Some statements are executed after `setDone()` and before the value of done is checked
 - In the worst case, `setDone()` is called right after the the done variable was checked
- **Delays** while waiting for a thread to terminate are **inevitable**
 - But it would be good if they could be minimized

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.36

Interrupting a thread

- When we arrange for thread to terminate, we:
 - Want it to **complete its blocking method** immediately
 - Don't wish to wait for the data (or ...) because the thread will exit
- Use `interrupt()` method of the `Thread` class to **interrupt** any **blocking method**

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.37

Effects of the interrupt method

- Causes blocked method to **throw an `InterruptedException`**
 - `sleep()`, `wait()`, `join()`, and methods to read I/O
- Sets a **flag** inside the thread object to indicate that the thread has been interrupted
 - Queried using `isInterrupted()`
 - Returns true if it was interrupted, even though it was not blocked

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.38

Stopping a thread: Using interrupts

```
public class RandomGen extends Thread {  
    public void run() {  
        while (!isInterrupted()) {  
            ...  
        }  
    }  
}  
  
randomGeneratorThread.interrupt();
```

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.39

The Runnable interface

- Allows **separation** of the **implementation** of the task **from the thread** used to run task

```
public interface Runnable {  
    public void run();  
}
```

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.40

Creation of a thread using the Runnable interface

- Construct the thread
 - Pass runnable object to the thread's constructor
- Start the thread
 - Instead of starting the runnable object

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.41

Creation of a thread using the Runnable interface

```
public class RandomGenerator implements Runnable {  
    public void run() { ... }  
}  
  
...  
generator = new RandomGenerator();  
Thread createdThread = new Thread(generator);  
createdThread.start();
```

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.42

When to use Runnable and Thread

- If you would like your class to inherit behavior from the Thread class
 - ▣ **Extend** Thread
- If your class needs to inherit from other classes
 - ▣ **Implement** Runnable

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.43

If you extend the Thread class?

- You **inherit behavior** and **methods** of the Thread class
 - ▣ The `interrupt()` method is part of the Thread class
 - ▣ You can `interrupt()` **if you extend**

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.44

Advantages of using the Runnable interface

- Java provides several classes that handle threading **for you**
 - ▣ Implement pooling, scheduling, or timing
 - ▣ These require the **Runnable** interface

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.45

But what if I still can't decide?

- Do a UML model of your application
- The object hierarchy tells you what you need:
 - ▣ If your task needs to subclass another class?
 - Use Runnable
 - ▣ If you need to use methods of Thread within your class?
 - Use Thread

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.46

Threads and Objects

- Instance of the Thread class is just an **object**
 - ▣ Can be passed to other methods
 - ▣ If a thread has a reference to another thread
 - It can invoke **any method** of that thread's object
- The Thread object is not the thread itself
 - ▣ It is the set of methods and data that **encapsulate** information about the thread

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.47

But what does this mean?

- You cannot look at the object source and **know which thread is:**
 - ▣ Executing its methods or examining its data
- You may wonder about which thread is running the code, but ...
 - ▣ There may be many possibilities

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.48

Determining the current thread

- Code within a thread object might want to see that code is being executed either:
 - By thread represented by the object or
 - By a completely different thread
- Retrieve reference to current thread
 - `Thread.currentThread()`
 - Static method

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.49

Checking which thread is executing the code

```
public class MyThread extends Thread {  
    public void run() {  
        if (Thread.currentThread() != this) {  
            throw new IllegalStateException(  
                "Run method called by incorrect thread -");  
        } /* end if */  
        ... Main logic  
    }  
}
```

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.50

Allowing a Runnable object to see if it has been interrupted

```
public class MyRunnable implements Runnable {  
    public void run() {  
        if (!Thread.currentThread().isInterrupted()) {  
            ... Main logic  
        }  
    }  
}
```

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.51

The contents of this slide-set are based on the following references

- Java Threads. Scott Oaks and Henry Wong. . 3rd Edition. O'Reilly Press. ISBN: 0-596-00782-5/978-0-596-00782-9. [Chapters 3, 4]
- Operating Systems Principles and Practice. Thomas Anderson and Michael Dahlin. 2nd Edition. ISBN: 978-0985673529. [Chapter 4]

February 1, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L6.52