

## CS 455: INTRODUCTION TO DISTRIBUTED SYSTEMS [THREADS]

Shrideep Pallickara  
 Computer Science  
 Colorado State University

February 6, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L7.1

## Frequently asked questions from the previous class survey

- Which is preferable: interrupting a thread or returning from method?
- Portion of heap: Are these preallocated? Equal in size?
- How is the threshold for deepest nesting level determined?
- Example of a checked exception? FileNotFoundException
- join(): why would you need this?
- Thread T1 has a method doSomething() that its run() never calls. When it is executing, does it mean some other thread T<sub>x</sub> called it?
- Threads and control of terminal when you launch it?
- Thread context switches vs Process context switches

February 6, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L7.2

## Topics covered in this lecture

- Data synchronization
- Synchronized blocks
- Lock fairness
- Wait-notify

February 6, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L7.3

## Heisenbugs

- Term coined by ACM Turing Award winner Jim Gray
  - Pun on the name of Werner Heisenberg
  - Act of observing a system, alters its state!
- Describes a particular class of bugs
  - Those that disappear or change behavior when you try to examine them
- Multithreaded programs are a common source of Heisenbugs

February 6, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L7.4

## What about regular bugs?

- Sometimes referred to as Bohr bugs
  - Deterministic
  - Generally much easier to diagnose

February 6, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L7.5

## Reasoning about interleaved access to shared state: Too much milk!

	Roommate 1's actions	Roommate 2's actions
3:00	Look in fridge; out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in fridge; out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home; put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home; put milk away
		<b>Oh no!</b>

February 6, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L7.6

## DATA SYNCHRONIZATION

February 6, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.7

## Why sharing data between threads is problematic

- **Race conditions**
- Threads attempt to access data more or less *simultaneously*
  - A thread may change the value of data that some other thread is operating on

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.8

## Example code with race condition

```
public class MyThread extends Thread {
    private byte[] values;
    private int position;

    public void
        modifyData(byte[] newValues, int newPosition) {
        ... Modify values and position
    }

    public void utilizeDataAndPerformFunction() {
        ... Use values and position
    }

    public void run() {
        ... Main logic
    }
}
```



February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.9

## In the previous snippet a race condition exists because ...

- The thread that calls modifyData() is **accessing the same data** as the thread that calls utilizeDataAndPerformFunction()
- utilizeDataAndPerformFunction() and modifyData() **are not atomic**
  - It is possible that values and position are changed *while they are being used*

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.10

## What is atomic?

- The code cannot be interrupted during its execution
  - Accomplished in hardware or *simulated* in software
- Code that cannot be found in an *intermediate state*

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.11

## Eliminating the race condition using the synchronized keyword

- If we declared both modifyData() and utilizeDataAndPerformFunction() as **synchronized**?
  - Only one thread gets to call *either* method at a time
    - Only one thread accesses data at a time
  - When one thread calls one of these methods, while another is executing one of them?
    - The second thread must *wait*

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.12


### Example code with no race conditions by using the synchronized keyword

```
public class MyThread extends Thread {
    private byte[] values;
    private int position;

    public void synchronized
        modifyData(byte[] newValues, int newPosition) {
        ... Modify values and position
    }

    public void synchronized
        utilizeDataAndPerformFunction() {
        ... Use values and position
    }

    public void run() {
        ... Main logic
    }
}
```



February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.13

### Revisiting the mutex lock

- **Mutually exclusive** lock
- If two threads try to grab a mutex?
  - Only one succeeds
- In Java every object has an associated **lock**

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.14

### When a method is declared synchronized ...

- The thread that wants to execute the method must **acquire** a lock
- Once the thread has acquired the lock?
  - It executes method and **releases** the lock
- When a method returns, the lock is released
  - Even if the return is because of an exception

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.15

### Locks and objects

- There is only **one lock per object**
- If two threads call synchronized methods of the same object?
  - Only one can execute immediately
    - The other has to wait until the lock is released

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University


L7.16

### Another code snippet to look at ...

```
public class MyThread extends Thread {
    private boolean done = false;

    public void run() {
        while (!done) {
            ... Main logic
        }
    }

    public void setDone(boolean isDone) {
        done = isDone;
    }
}
```



February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.17

### Can't we just synchronize the two methods as we did previously?

- If we synchronized both `run()` and `setDone()` ?
  - `setDone()` would never execute!
- The `run()` method does not exit until the done flag is set
  - But the done flag cannot be set because `setDone()` cannot execute till `run()` completes
- Uh oh ...

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.18

### The problem stems from the scope of the lock

- **Scope of a lock**
  - Period between grabbing and releasing a lock
- Scope of the `run()` method is too large!
  - Lock is grabbed and never released
- We will look at techniques to *shrink the scope* of the lock
- But let's look at another solution for now

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.19

### Let's look at operations performed on the data item (done)

- The `setDone()` method stores a value into the flag
- The `run()` method reads the value
- In our previous example:
  - Threads were accessing *multiple* pieces of data
  - No way to update multiple data items *atomically* without the synchronized keyword

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.20

### But Java specifies that the loading and storing of variables is atomic

- Except for long and double variables
- The `setDone()` should be atomic
  - The `run()` method has only one read operation of the data item
- The race condition should not exist
  - But why is it there?

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.21

### Threads are allowed to hold values of variables in registers

- When one thread changes the value of the variable?
  - Another thread *may not see* the changed variable
- This is particularly true in loops controlled by a variable
  - Looping thread **loads value of variable in register** and *does not notice* when value is changed by another thread

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.22

### Two approaches to solving this

- Providing setter and getter methods for variable and using the synchronized keyword
  - *When lock is acquired*, temporary values stored in registers are *flushed* to main memory
- The **volatile** keyword
  - Much cleaner solution

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.23

### If a variable is marked as volatile

- Every time it is used?
  - Must be read from main memory
- Every time it is written?
  - Must be written to main memory
- Load and store operations are **atomic**
  - Even for long and double variables

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.24

### Some more about volatile variables

- Prior to JDK 1.2 variables were always read from main memory
  - ▣ Using volatile variables was moot
- Subsequent versions introduced memory models and optimizations

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.25

### Synchronization and the volatile keyword

- Can be used *only* when operations use a **single load and store**
  - ▣ Operations like ++, --?
    - Load-change-store ...
- The volatile keyword forces the JVM to not make temporary copies of a variable
- Declaring an array volatile?
  - ▣ The reference becomes volatile
  - ▣ The individual elements are not volatile

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.26

### SYNCHRONIZED METHODS & LOCKS

February 6, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.27

### Synchronizing methods

- **Not possible** to execute the same method in one thread while ...
  - ▣ Method is running in another thread
- If two different synchronized methods in an object are called?
  - ▣ They both require the lock of the same object
- Two or more synchronized methods of the same object *can never run in parallel* in separate threads

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.28

### A lock is based on a specific instance of an object

- Not on a particular method or class
- Suppose we have 2 objects: objectA and objectB with synchronized methods modifyData() and utilizeData()
- One thread can execute objectA.modifyData() while another executes objectB.utilizeData() *in parallel*
  - ▣ Two different locks are grabbed by two different threads
  - ▣ No need for threads to wait for each other

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.29

### How does a synchronized method behave in conjunction with an unsynchronized one?

- Synchronized methods try to grab the object lock
  - ▣ Only 1 synchronized method in a object can run at a time ... *provides data protection*
- Unsynchronized methods
  - ▣ Don't grab the object lock
  - ▣ Can *execute at any time ... by any thread*
    - Regardless of whether a synchronized method is running

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.30

For a given object, at any time ...

- **Any number** of *unsynchronized methods* may be executing
- But only **1 synchronized method** can execute

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.31

Synchronizing static methods

- A lock can be obtained for each class
  - **The class lock**
- The class lock is the *object lock* of the **Class object** that models the class
  - There is only 1 **Class** object per class
  - Allows us to achieve synchronization for static methods

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.32

Object locks and class locks

- Are **not operationally related**
- The class lock can be grabbed and released independently of the object lock
- If a non-static synchronized method calls a static synchronized method?
  - It acquires both locks

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.33

## EXPLICIT LOCKING

February 6, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.34

The synchronized keyword

- *Serializes accesses* to synchronized methods in an object
- Not suitable for *controlling lock scope* in certain situations
- Can be *too primitive* in some cases

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.35

Many synchronization schemes in J2SE 5.0 onwards implement the `Lock` interface

- Two important methods
  - `lock()` and `unlock()`
- Similar to using the synchronized keyword
  - Call `lock()` at the start of the method
  - Call `unlock()` at the end of the method
- Difference: we have an *actual object* that **represents** the lock
  - Store, pass around, or discard

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.36

## Semantics of the using Lock

- If another thread *owns* the lock
  - Thread that attempts to acquire the lock must wait until the other thread calls `unlock()`
- Once the waiting thread acquires the lock, it returns from the `lock()` method

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.37

## Using the Lock interface

```
public class DataOperator {  
    private Lock dataLock = new ReentrantLock();  
    public void  
    modifyData(byte[] newValues, int newPosition) {  
        try {  
            dataLock.lock();  
            ... Modify values and position  
        } finally {  
            dataLock.unlock();  
        }  
    }  
  
    public void utilizeDataAndPerformFunction() {  
        try {  
            dataLock.lock();  
            ... Use values and position  
        } finally {  
            dataLock.unlock();  
        }  
    }  
}
```

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.38

## Advantages of using the Lock interface

- Grab and release locks *whenever* we want
- Now possible for **two objects to share the same lock**
  - Lock is no longer attached to the object whose method is being called
- Can be *attached to data, groups of data*, etc.
  - Not objects containing the executing methods

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.39

## Advantages of explicit locking

- We can move them anywhere to **adjust lock scope**
  - Can span from a line of code to a scope that encompasses multiple methods and objects
- Lock at scope *specific to problem*
  - Not just the object

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.40

## SYNCHRONIZED BLOCKS

February 6, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.41

## Much of what we accomplish with the Lock we can do so with the synchronized keyword

```
public class DataOperator {  
    public void  
    modifyData(byte[] newValues, int newPosition) {  
        synchronized(this) {  
            ... Modify values and position  
        }  
    }  
  
    public void utilizeDataAndPerformFunction() {  
        synchronized(this) {  
            ... Use values and position  
        }  
    }  
}
```

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.42

## Synchronized methods vs. Synchronized Blocks

- Possible to use only the **synchronized block** mechanism to synchronize whole method
- You decide when it's best to synchronize a block of code or the whole method
- RULE: **Establish as small a lock scope as possible**

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.43

## The Lock interface [java.util.concurrent.locks]

```
public interface Lock {  
    public void lock();  
    public void lockInterruptibly() throws InterruptedException;  
    public boolean tryLock();  
    public boolean tryLock(long time, TimeUnit unit) throws InterruptedException;  
    public void unlock();  
    public Condition newCondition();  
}
```

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.44

## Lock Fairness

- ReentrantLock allows locks to be granted **fairly**
  - Locks are granted as close to arrival order as possible
  - Prevents **lock starvation** from happening
- Possibilities for granting locks
  - ① First-come-first-served
  - ② Allows servicing the maximum number of requests
  - ③ Do what's best for the platform

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.45

## THREAD NOTIFICATIONS

February 6, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.46

## Objects and communications

- Every object has a lock
- Every object also includes mechanisms that allow it to be a **waiting area**
  - Allows **communication** between threads

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.47

## Conditions

- One thread needs a **condition** to exist
  - Assumes another thread will **create** that condition
- When another thread creates the condition?
  - It **notifies** the first thread that has been **waiting** for that condition

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.48



## wait(), notify() and the Object class

```
public class Object {  
    public void wait();  
    public void wait(long timeout);  
    public void notify();  
}
```

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.49

## wait(), notify() and the Object class

- Wait-and-notify mechanisms are available for every object
  - Accomplished by **method invocations**
- Synchronized mechanism is handled by using a **keyword**

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.50

## Wait-and-notify relate to synchronization, but ...

- It is more of a **communications mechanism**
- Allows one thread to communicate to another that a **condition** has occurred
  - Does not specify **what** that specific condition is

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.51

## Can wait-and-notify replace the synchronized mechanism?

- No
- Does not** solve the race condition that the synchronized mechanism solves
- Must be used in **conjunction** with the synchronized lock
  - Prevents race condition that exists in the wait-notify mechanism itself

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.52

## A code snippet that uses wait-notify to control the execution of the thread

```
public class Tester implements Runnable {  
    private boolean done = true;  
  
    public synchronized run() {  
        while (true) {  
            if (done) wait();  
            else { ... Logic ... wait(100);}  
        }  
    }  
  
    public synchronized void setDone(boolean b) {  
        done = b;  
        if (!done) notify();  
    }  
}
```

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.53

## About the wait() method

- When wait() executes, the synchronization lock is **released**
  - By the JVM
- When a notification is received?
  - The thread needs to **reacquire** the synchronization lock before returning from wait()

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.54

### Integration of wait-notify and synchronization

- **Tightly integrated** with the synchronization lock
  - Feature **not directly available to us**
  - Not possible to implement this: native method
- This is typical of approach in other libraries
  - **Condition variables** for Solaris and POSIX threads require that a mutex lock be held

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.55

### Details of the race condition in the wait-notify mechanism

- The first thread **tests the condition** and confirms that it must wait
- The second thread **sets the condition**
- The second thread calls `notify()`
  - This **goes unheard** because the first thread is not yet waiting
- The first thread calls `wait()`

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.56

### How does the potential race condition get resolved?

- To call `wait()` or `notify()`
  - Obtain lock for the object on which this is being invoked
- It seems as if the lock has been held for the entire `wait()` invocation, but ...
  - ① `wait()` **releases lock prior to waiting**
  - ② **Reacquires the lock just before returning** from `wait()`

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.57

### Is there a race condition during the time `wait()` releases and reacquires the lock?

- `wait()` is **tightly integrated** with the lock mechanism
- Object lock is **not freed until** the waiting thread is in a **state in which it can receive notifications**
  - System prevents race conditions from occurring here

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.58

### If a thread receives a notification is it guaranteed that condition is set?

- No
- **Prior** to calling `wait()`, **test condition** while holding lock
- Upon **returning** from `wait()` **retest** condition to see if you should `wait()` again

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.59

### What if `notify()` is called and no thread is waiting?

- Wait-and-notify mechanism has no knowledge about the condition about which it notifies
- If `notify()` is called when no other thread is waiting?
  - The notification is lost

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.60

### What happens when more than 1 thread is waiting for a notification?

- Language specification does not define which thread gets the notification
  - Based on JVM implementation, scheduling and timing issues
- *No way to determine* which thread will get the notification

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.61

### notifyAll()

- All threads that are waiting on an object are notified
- When threads receive this, they must work out
  - ① Which thread should continue
  - ② Which thread(s) should call `wait()` again
    - All threads wake up, but they **still have to reacquire the object lock**
    - Must wait for the lock to be freed

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.62

### The contents of this slide-set are based on the following references

- *Java Threads*. Scott Oaks and Henry Wong. 3rd Edition. O'Reilly Press. ISBN: 0-596-00782-5/978-0-596-00782-9. [Chapter 3, 4]
- *Operating Systems Principles and Practice*. Thomas Anderson and Michael Dahlin. 2nd Edition. ISBN: 978-0985673529. [Chapter 5]

February 6, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L7.63