

CS 455: INTRODUCTION TO DISTRIBUTED SYSTEMS [THREAD SAFETY]

Shrideep Pallickara
Computer Science
Colorado State University

February 8, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.1

Frequently asked questions from the previous class survey

- synchronize in interfaces?
- Multiple inheritance and synchronized
 - Java supports *multiple inheritance of type* NOT implementation
- If a Thread A calls a synchronized method in Thread B, does it acquire the object lock for B ... what would happen to A?
- Any other way to bypass using volatile (without wait-notify or synchronized)
- wait-notify and the Lock interface
- When using synchronized block in a synchronized method which locks would you acquire?
- What if an unsynchronized method calls a synchronized method? Can other threads execute in the unsynchronized method?
- Are locks inherited?

February 8, 2018

Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.2

Topics covered in this lecture

- Thread safety
- Compound actions
- Reentrancy
- Sharing objects and confinement

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.3

A code snippet that uses wait-notify to control the execution of the thread

```
public class Tester implements Runnable {  
    private boolean done = true;  
  
    public synchronized run() {  
        while (true) {  
            if (done) wait();  
            else { ... Logic ... wait(100);}  
        }  
    }  
  
    public synchronized void setDone(boolean b) {  
        done = b;  
        if (!done) notify();  
    }  
}
```

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.4

Details of the race condition in the wait-notify mechanism

- The first thread *tests the condition* and confirms that it must wait
- The second thread *sets the condition*
- The second thread calls `notify()`
 - This *goes unheard* because the first thread is not yet waiting
- The first thread calls `wait()`

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.5

How does the potential race condition get resolved?

- To call `wait()` or `notify()`
 - Obtain lock for the object on which this is being invoked
- It seems as if the lock has been held for the entire `wait()` invocation, but ...
 - ① `wait()` releases lock prior to waiting
 - ② Reacquires the lock just before returning from `wait()`

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.6

Is there a race condition during the time `wait()` releases and reacquires the lock?

- `wait()` is **tightly integrated** with the lock mechanism
- Object lock is **not freed until** the waiting thread is in a **state in which it can receive notifications**
 - System prevents race conditions from occurring here

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.7

If a thread receives a notification is it guaranteed that condition is set?

- No
- **Prior** to calling `wait()`, **test condition** while holding lock
- Upon **returning** from `wait()` **retest** condition to see if you should `wait()` again

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.8

What if `notify()` is called and no thread is waiting?

- Wait-and-notify mechanism has no knowledge about the condition about which it notifies
- If `notify()` is called when no other thread is waiting?
 - The notification is lost

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.9

What happens when more than 1 thread is waiting for a notification?

- Language specification does not define which thread gets the notification
 - Based on JVM implementation, scheduling and timing issues
- **No way to determine** which thread will get the notification

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.10

`notifyAll()`

- All threads that are waiting on an object are notified
- When threads receive this, they must work out
 - ① Which thread should continue
 - ② Which thread(s) should call `wait()` again
 - All threads wake up, but they **still have to reacquire the object lock**
 - Must wait for the lock to be freed

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.11

Threads and locks

- **Locks are held by threads**
 - A thread can hold multiple locks
 - Any thread that tries to obtains these locks? Placed into a wait state
 - If the thread deadlocks? It results in all locks that it holds becoming unavailable to other threads
- If a lock is held by some other thread?
 - The thread **must wait** for it to be free: **There is no preemption of locks!**
 - If the lock is unavailable (or held by a deadlocked thread) it blocks all the waiting threads

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.12

THREAD SAFETY

February 8, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.13

Race conditions

- Getting the right answer depends on lucky timing
 - E.g. check-then-act: When stale observations are used to make a decision on what to do next
- Real world example
 - 2 friends trying to meet up for coffee on campus without specifying which of the 2 locations

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.14

Racing and synchronization

[1/3]

- Purpose of **synchronization**?
 - Prevent **race conditions** that can cause data to be found in either an inconsistent or intermediate state
- Threads are not allowed to race during sections of code protected by synchronization
 - But this does not mean outcome or order of execution of threads is deterministic
 - Threads may be racing prior to the synchronized section of code

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.15

Racing and synchronization

[2/3]

- If threads are waiting on the same lock
 - The order in which the synchronized code is executed is determined by order in which lock is granted
 - Which is platform-specific and non-deterministic

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.16

Racing and synchronization

[3/3]

- Not all races should be avoided
 - This is a subtle but important point: If you do this ... every thing is serialized
 - **Only race-conditions within thread-unsafe sections of the code** are considered a problem
 - ① Synchronize code that prevents race condition
 - ② Design code that is thread-safe without the need for synchronization (or minimal synchronization)

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.17

CONCURRENT PROGRAMMING

February 8, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.18

Concurrent programming

- Concurrent programs require the **correct use** of threads and locks
- But these are just **mechanisms**

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.19

Object State

- Includes its **data**
 - Stored in instance variables or static fields
 - Fields from dependent objects
 - HashMap's state also depends on Map.Entry<K, V> objects
- Encompasses any data that can affect its **externally visible** behavior

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.20

The crux of developing thread safe programs

- Managing access to **state**
 - In particular **shared, mutable state**
- Shared
 - Variables could be accessed by multiple threads
- Mutable
 - Variable's values change over its lifetime
- Thread-safety
 - **Protecting data from uncontrolled concurrent access**

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.21

When to coordinate accesses

- Whenever more than one thread accesses a state variable, and one of them **might write** to it?
 - They must all coordinate their access to it
- Avoid temptation to think that there are special situations when you can disregard this

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.22

When should an object be thread-safe?

- Will it be accessed from multiple threads?
- The key here is **how** the object is **used**
 - Not **what** it **does**

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.23

How to make an object thread-safe

- Use **synchronization** to **coordinate** access to mutable state
- Failure to do this?
 - Data corruptions
 - Problems that manifest themselves in myriad forms

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.24

Mechanisms for synchronization in Java

- One way to achieve this is via the **synchronized** keyword
 - Exclusive locking
- Other approaches include:
 - **volatile** variables
 - Explicit **locks**
 - **Atomic** variables

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.25

Programs that omit synchronizations

- Might work for some time
 - But it **will break** at some point
- Far easier to design a class to be thread-safe **from the start**
 - Retrofitting it to be thread-safe is extremely hard

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.26

Thread-safety: Encapsulate your state

- Fewer code should have access to a particular variable
 - Easier to reason about **conditions** under which it might be accessed
- **DON'T:**
 - Store state in public fields
 - Publish reference to an **internal** object



February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.27

Fixing access to mutable state variables from multiple threads

- **Don't share** state variables across threads
- Make state variables **immutable**
- Use **synchronization** to coordinate access to the state variable

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.28

Correctness of classes

- Class conforms to **specification**
- **Invariants** constrain object's state
- **Post conditions** describe the effects of operations

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.29

A Thread-safe class

- **Behaves correctly** when accessed from multiple threads
- Regardless of **scheduling or interleaving** of execution of those threads
 - By the runtime environment
- No additional synchronization or coordination by the calling code

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.30

Really?

- Thread safe classes encapsulate *any needed* synchronization
- Clients should not have to provide their own

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.31

Stateless objects are thread-safe

```
public class StatelessClass implements Servlet {  
    public void factorizer(ServletRequest req,  
        ServletResponse resp) {  
        BigInteger i = extractFromReq(req);  
        BigInteger[] factors = factorize(i);  
        encodeIntoResponse(resp, factors);  
    }  
}
```

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.32

Stateless objects are always thread-safe

- **Transient state** for a particular computation exists solely in *local variables*
 - Stored on the thread's stack
 - Accessible only to the executing thread
- One thread cannot influence the result of another
 - The threads have no shared state

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.33

Atomicity

- Let's look at two operations **A** and **B**
- From the perspective of thread executing **A**
- When another thread executes **B**
 - Either all of **B** has executed or none of it has
- Operation **A** and **B** are **atomic with respect to each other**

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.34

Initializing Objects

```
public class LazyInitialization {  
    private ExpensiveObject instance = null;  
    public ExpensiveObject getInstance() {  
        if (instance == null) {  
            instance = new ExpensiveObject();  
        }  
        return instance;  
    }  
}
```



February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.35

Thread-safe initialization

```
public class Singleton {  
    private static final Singleton instance = new Singleton();  
    // Private constructor prevents instantiation from other  
    // classes  
    private Singleton() { }  
    public static Singleton getInstance() {  
        return instance;  
    }  
}
```



February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.36

The final keyword

- You cannot extend a final class
 - E.g. `java.lang.String`
- You cannot override a final method
- You can only initialize a final variable **once**
 - Either via an initializer or an assignment statement

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.37

Blank final instance variable of a class

- Must be assigned **within every constructor** of the class
- Attempting to set it outside the constructor will result in a compilation error
- The value of a final variable is not necessarily known at compile time

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.38

Atomicity with compound operations

```
public class CountingFactorizer {
    private long count = 0;

    public long getCount() {return count;}

    public void factorizer(int i) {
        int[] factors = factor(i);
        count++;
    }
}
```



February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.39

Atomicity with compound operations

```
public class CountingFactorizer {
    private final AtomicLong count = new AtomicLong(0);

    public long getCount() {return count;}

    public void factorizer(int i) {
        int[] factors = factor(i);
        count.incrementAndGet();
    }
}
```



February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.40

Compound actions & thread-safety

- Compound actions
 - Check-then-act
 - Read-modify-write
- Must be executed atomically for thread-safety

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.41

LOCKS & REENTRANCY

February 8, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.42

Reentrancy

- When thread requests lock held by another thread?
 - ▣ Requesting thread blocks
- If a thread attempts to acquire a lock it already holds?
 - ▣ Succeeds
- Locks are acquired on a **per-thread** rather than on a per-invocation basis

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.43

How reentrancy works

[1/2]

- For each lock two items are maintained
 - ▣ Acquisition count
 - ▣ Owning thread
- When the count is zero?
 - ▣ Lock is free
- If a thread acquires lock for the first time?
 - ▣ Count is one

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.44

How reentrancy works

[2/2]

- If owning thread acquires lock again, count is incremented
- When owning thread exits synchronized block, count is decremented
 - ▣ If it is zero Lock is released

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.45

Does this result in a deadlock?

```
public class Widget {
    public synchronized doSomething() {
        ...
    }
}

public class LoggingWidget extends Widget {
    public synchronized void doSomething() {
        System.out.println(toString()+"Calling doSomething()");
        super.doSomething();
    }
}
```

No! Intrinsic locks are reentrant



February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.46

GUARDING STATE WITH LOCKS

February 8, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.47

Guarding state with locks

- A **mutable, shared** variable that may be accessed by multiple threads must be guarded by the **same lock**
- For every **invariant** that involves more than one variable?
 - ▣ **All variables** must be guarded by the **same lock**

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.48

Watch for indiscriminate use of synchronization

- Every method in Vector is synchronized
- But this does not render compound actions on Vector atomic

```
if (!vector.contains(element)) {  
    vector.add(element);  
}
```

- Snippet has *race condition* even though add and contains are atomic
- **Additional locking needed for compound actions**

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.49

Pitfalls of over synchronization

- Number of simultaneous invocations?
 - Not limited by processor resources, but is limited by the application structure
 - **Poor concurrency**

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.50

Antidote for poor concurrency

- Control the **scope** of the lock
 - Too large: Invocations become sequential
 - Don't make it too small either
 - Operations that are atomic should not be in synchronized block

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.51

The contents of this slide-set are based on the following references

- *Java Threads*. Scott Oaks and Henry Wong. 3rd Edition. O'Reilly Press. ISBN: 0-596-00782-5/978-0-596-00782-9. [Chapter 3, 4]
- *Java Concurrency in Practice*. Brian Goetz, Tim Peierls, Joshua Bloch, Joseph Bowbeer, David Holmes, and Doug Lea. Addison-Wesley Professional. ISBN: 0321349601/978-0321349606. [Chapters 1, 2, 3 and 4]

February 8, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L8.52