

CS 455: INTRODUCTION TO DISTRIBUTED SYSTEMS [THREAD SAFETY]

Putting the brakes, on impending code breaks

Let a reference escape, have you?

Misbehave, your code will, out of the blue

Get out, you will, of this bind

If, your objects, you have confined

Shrideep Pallickara
Computer Science
Colorado State University

February 13, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.1

Frequently asked questions from the previous class survey

- What if A calls B, which is atomic, and B fails?
- Externally visible behavior of an object?
- Is a class with private final objects thread safe?

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.2

Topics covered in this lecture

- Sharing & Composing Objects
- Making a class thread-safe
- Multivariable invariants and thread-safety
- Adding functionality to a thread-safe class

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.3

Pitfalls of over synchronization

- Number of simultaneous invocations?
 - Not limited by processor resources, but is limited by the application structure
 - **Poor concurrency**

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.4

Antidote for poor concurrency

- Control the **scope** of the lock
 - Too large: Invocations become sequential
 - Don't make it too small either
 - Operations that are atomic should not be in a synchronized block

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.5

SHARING OBJECTS

February 13, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.6

What we will be looking at

- Techniques for sharing and publishing objects
 - Safe access from multiple threads
- Together with synchronization, sharing objects lays foundation for thread-safe classes

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.7

Synchronization

- What we have seen so far:
 - Atomicity and demarcating *critical sections*
- But it is also about **memory visibility**
 - We *prevent* one thread from modifying object state while another is using it
 - When *state of an object is modified*, other thread can *see* the changes that were made

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.8

Publication and Escape

- Publishing an object
 - Makes it available *outside* current scope
 - Storing a reference to it, returning from a non-private method, passing it as argument to another method
- **Escape**
 - An object that is published when it *should not* have been

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.9

Pitfalls in publication

- Publishing internal state variables
 - Makes it *difficult* to preserve invariants
- Publishing objects before they are constructed
 - Compromises thread-safety

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.10

Most blatant form of publication

- Storing a reference in a public static field

```
public static Set<Secret> knownSecrets;  
  
public void initialize() {  
    knownSecrets = new HashSet<Secret>();  
}
```



- If you add a Secret to knownSecrets?
- You also end up publishing that Secret

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.11

Allowing internal mutable state to escape

```
public class PublishingState {  
    private String[] states = new String[] {  
        "AK", "AL", ...  
    };  
  
    public String[] getStates() {return states;}  
}
```



- states has *escaped* its intended scope
 - What should have been private is now public
- **Any caller can modify its contents**

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.12

Another way to publish internal state

```
public class ThisEscape {  
    public ThisEscape(EventSource source) {  
        source.registerListener(  
            new EventListener() {  
                public void onEvent(Event e) {  
                    doSomething(e);  
                }  
            });  
    }  
}
```



- When EventListener is published, it publishes the enclosing ThisEscape instance
- **Inner class instances contain hidden reference to enclosing instance**

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.13

Safe construction practices

- An object is in a predictable, consistent state *only after its constructor returns*
- Publishing an object within its constructor?
 - You are publishing an incompletely constructed object
 - Even if you are doing so in the last line of the constructor
- **RULE: Don't allow `this` to escape during construction**

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.14

A common mistake is to start a thread from a constructor

- When an object creates a thread in its constructor
 - Almost always shares its `this` reference with the new thread
 - Explicitly: Passing it to the constructor
 - Implicitly: The Thread or Runnable is an inner class of the owning object
- Nothing wrong with creating a thread in a constructor
 - Just don't start the Thread
 - Expose an `initialize()` method

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.15

THREAD CONFINEMENT

February 13, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.16

Thread confinement

- Accessing shared, mutable data requires synchronization
 - Avoid this by *not sharing*
- If data is only accessed from a single thread?
 - No synchronization is needed
- When an object is **confined** to a thread?
 - Usage is **thread-safe even if the object is not**

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.17

Thread confinement

- Language has no means of confining an object to a thread
- Thread confinement is an element of **program's design**
 - Enforced by implementation
- Language and core libraries provide mechanisms to help with this
 - Local variables and the ThreadLocal class

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.18

Stack confinement

- Object can only be reached through local variables
- Local variables are **intrinsically confined** to the executing thread
 - Exist on executing thread's stack
 - Not accessible to other threads

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.19

Thread confinement of reference variables

```
public int loadTheArk() {
    SortedSet<Animal> animals;

    // animals confined to method don't let
    // them escape

    return numPairs;
}
```

If you were to publish a reference to animals,
stack confinement would be violated

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.20

ThreadLocal

- Allows you to associate a per-thread value with a value-holding object
- Provides set and get accessor methods
 - Maintains a separate copy of value for each thread that uses it
 - get returns the most recent value passed to set
 - From the currently executing thread

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.21

Using ThreadLocal for thread confinement

```
private static ThreadLocal<Connection> connectionHolder
= new ThreadLocal<Connection>() {
    public Connection initialValue() {
        return DriverManager.getConnection(DB_URL);
    }
};

public static Connection getConnection() {
    return connectionHolder.get();
}
```

Each thread will have its own connection

When thread calls ThreadLocal.get for the first time?
initialValue() provides the initial value

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.22

Common use of ThreadLocal

- Used when a frequently used operation requires a temporary object
 - Wish to avoid reallocating temporary object on each invocation
- Integer.toString()
 - Before 5.0 used ThreadLocal to store a 12-byte buffer for formatting result

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.23

IMMUTABLE OBJECTS

February 13, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.24

Immutable objects

- State cannot be modified after construction
- All its fields are `final`
- Properly constructed
 - The `this` reference does not escape during construction

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.25

Immutable objects

```
public final class ThreeStooges {  
    private final Set<String> stooges = new HashSet<String>();  
  
    public ThreeStooges() {  
        stooges.add("Moe");  
        stooges.add("Larry");  
        stooges.add("Curly");  
    }  
  
    public boolean isStooge() {return stooges.contains(name);}  
}
```

Design makes it impossible to modify after construction

The `stooges` reference is `final`

All object state reached through a `final` field

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.26

Safe publication of objects

- Storing reference to an object into a public field is **not enough** to publish that object safely

```
public Holder holder;  
  
public void initialize() {  
    holder = new Holder(42);  
}
```



Holder could appear to be in an inconsistent state

Even though invariants may have been established by constructor

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.27

Class at risk of failure if not published properly

```
public class Holder {  
    private int n;  
  
    public Holder(int n) {this.n == n}  
  
    public void assertSanity() {  
        if (n != n) {  
            throw new AssertionError("Statement is false");  
        }  
    }  
}
```



Thread may see a stale value first time it reads the field and
an up-to-date value the next time

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.28

COMPOSING OBJECTS

February 13, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.29

Composing Objects

- We don't want to have to analyze *each memory access* to ensure program is thread-safe
- We wish to take thread-safe components and **compose** them into larger components or programs

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.30

Basic elements of designing a thread-safe class

- Identify **variables** that **form** the object's state
- Identify **invariants** that **constrain** the state variables
- Establish a **policy** for managing **concurrent access** to the object's state

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.31

Synchronization policy

- Defines how object **coordinates access** to its state
 - Without violating its invariants or post-conditions
- Specifies a **combination** of:
 - Immutability
 - Thread confinement
 - Locking

} To maintain
Thread Safety

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.32

Looking at a counter

```
public final class Counter {
    private long value=0;

    public synchronized long getValue() {
        return value;
    }

    public synchronized long increment() {
        if (value == Long.MAX_VALUE) {
            throw new IllegalStateException("Counter Overflow");
        }
        value++;
        return value;
    }
}
```



February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.33

Making a class thread-safe

- Ensure that invariants hold under concurrent access
 - We need to **reason** about state
- Object and variables have **state space**
 - **Range** of possible states
 - **Keep this small** so that it is easier to reason about

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.34

Classes have invariants that tag certain states as valid or invalid

- Looking back at our **Counter** example
- The value field is a long
- The state space ranges from Long.MIN_VALUE to Long.MAX_VALUE
- The class places constraints on value
 - Negative values are not allowed

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.35

Operations may have post conditions that tag state transitions as invalid

- Looking back at our **Counter** example
- If the current state of Counter is 17
 - The **only** valid next state is 18
 - When the next state is **derived from the current state?**
 - **Compound action**
- Not all operations impose state transition constraints
 - For e.g. if a variable tracks current temperature? Previous state doesn't impact current state

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.36

Constraints and synchronization requirements

- If certain states are invalid?
 - Underlying state variables should be **encapsulated**
 - If not, client code can put it in an *inconsistent* state
- If an operation has invalid state transitions?
 - It must be made **atomic**

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.37

Looking at a case where invariants constrain multiple state variables

```
public class NumberRange {
    private final AtomicInteger lower = new AtomicInteger(0);
    private final AtomicInteger upper = new AtomicInteger(0);

    public void setLower(int i) {
        if (i > upper.get())
            throw IllegalArgumentException("lower > upper!");
        lower.set(i);
    }

    public void setUpper(int i) {
        if (i < lower.get())
            throw IllegalArgumentException("upper < lower!");
        upper.set(i);
    }

    public boolean isInRange(int i) {
        return (i >= lower.get() && i <= upper.get());
    }
}
```



February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.38

Problems with NumberRange

- Does not preserve invariant that constrains lower and upper
- The methods setLower and setUpper *attempt* this preservation
 - But they do so poorly!
 - They are *check-then-act* sequences that use *insufficient locking* that precludes atomicity

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.39

Problems with NumberRange

- If the number range (0, 10) holds
- One thread calls setLower(5) while another calls setUpper(4)
- With unlucky timing?
 - Both calls will pass checks in the setters
 - Both modifications will be applied
- Range is now (5, 4) ... an invalid state
- AtomicInteger is thread-safe, the composite class is not

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.40

Multivariable invariants

- Related variables must be *fetched or updated* in an **atomic** operation
- Don't:
 - Update one
 - Release and reacquire lock, and ...
 - Then update others
- The lock that guards the variables
 - Must be **held for the duration of any operation** that accesses them

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.41

State-dependent operations

- Objects may have state-based **pre-conditions**
 - E.g., cannot remove item from an empty queue
- In a single-threaded program
 - Operations simply fail
- In a concurrent program
 - Precondition may be *true later* because of the actions of another thread

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.42

State dependent operations: Mechanisms

- `wait()/notify()`
 - ▣ Supported by the JVM and closely tied with intrinsic locking
- Other possibilities
 - ▣ Use classes such as blocking queues or semaphores

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.43

STATE OWNERSHIP

February 13, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.44

State ownership

- Defining which variables form an object's state
 - ▣ We wish to consider only that which the object owns
- Ownership
 - ▣ Not explicitly specified in the language
 - ▣ **Element of program design**

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.45

State ownership: Encapsulation and ownership go together

- Object encapsulates the state it owns
 - ▣ Owns the state it encapsulates
- Owner gets to decide on the **locking protocol**
- If you publish a reference to a mutable object?
 - ▣ You no longer have exclusive control

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.46

Instance confinement

- Object may not be thread-safe
 - ▣ But we could still use it in a thread-safe fashion
- Ensure that:
 - ▣ It is accessed by only one thread
 - ▣ All accesses guarded by a lock

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.47

Confinement and locking working together

```
public class PersonSet {  
    private final Set<Person> mySet = new HashSet<Person>();  
  
    public synchronized void addPerson(Person p) {  
        mySet.add(p);  
    }  
  
    public synchronized boolean containsPerson(Person p) {  
        return mySet.contains(p);  
    }  
}
```

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.48

Looking at our previous example

- State of `PersonSet` managed by `HashSet`, which is not thread-safe
- But `mySet` is
 - ▢ Private
 - ▢ Not allowed to escape
 - ▢ Confined to `PersonSet`

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.49

But we have made no assumptions about `Person`

- If it is mutable, additional synchronization is needed
 - ▢ When accessing `Person` from `PersonSet`
- Reliable way to achieve this?
 - ▢ Make `Person` thread-safe
- Less-reliable way?
 - ▢ Guard `Person` objects with a lock
 - ▢ Ensure that clients follow protocol of acquiring appropriate lock, before accessing `Person`

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.50

Instance confinement is the easiest way to build thread-safe classes

- Class that confines its state can be analyzed for thread-safety
 - ▢ Without having to examine the whole program

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.51

GUARDING STATE WITH PRIVATE LOCKS

February 13, 2018

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.52

Guarding state with a private lock

```
public class PrivateLock {
    private final Object myLock = new Object();
    private Widget widget; //guarded by myLock

    public void someMethod() {
        synchronized(myLock) {
            //Access and modify the state of the widget
        }
    }
}
```

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.53

Why guard state with a private lock?

- Doing so encapsulates the lock
 - ▢ **Client code cannot acquire it!**
- Publicly accessible lock allows client code to participate in its synchronization policy
 - ▢ Correctly or incorrectly
- Clients that improperly acquire an object's lock cause **liveness** issues
- Verifying correctness with public locks requires examining the entire program not just a class

February 13, 2018
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]
Dept. Of Computer Science, Colorado State University

L9.54

The contents of this slide-set are based on the following references

- *Java Concurrency in Practice*. Brian Goetz, Tim Peierls, Joshua Bloch, Joseph Bowbeer, David Holmes, and Doug Lea. Addison-Wesley Professional. ISBN: 0321349601/978-0321349606. [Chapters 3 and 4]