

## CS 455: INTRODUCTION TO DISTRIBUTED SYSTEMS [THREAD SAFETY & MAPREDUCE]

### Are you set on reinventing the wheel?

Shunning libraries and frameworks, are you, despite the peril?  
Emerge scathed, from arduous projects, you will

Survived, these have, the scrutiny of a thousand probing eyes  
Abrogating your choice, is what this implies

Shrideep Pallickara  
Computer Science  
Colorado State University

February 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.1

## Frequently asked questions from the previous class survey

- ConcurrentHashMap
  - Does the lock operate over a consecutive space?
  - During resize operations can elements be added/removed?
- Latches:
  - Why not use a counter object, that is guarded by synchronous methods?

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.2

## Topics covered in this lecture

- Thread safety wrap-up
  - Synchronizers and summary
- Map Reduce

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.3

## SYNCHRONIZERS

February 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.4

## Semaphores

- Counting semaphores control the **number of activities** that can:
  - Access a certain resource
  - Perform a given action
- Used to implement resource pools or impose bounds on a collection

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.5

## Semaphores

- Manage a set of virtual **permits**
  - Initial number passed to the constructor
- Activities **acquire** and **release** permits
- If **no permits** are available?
  - **acquire blocks** until one is available
- The **release** method returns a permit to the semaphore

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.6

## Semaphores are useful for implementing resource pools

- Block if the pool is empty
  - Unblock if the pool is non-empty
- Initialize a semaphore to the **pool size**
- acquire a permit before trying to fetch a resource from pool
- release the permit after putting the resource back in pool
- acquire **blocks** until the pool is non-empty

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.7

## Binary semaphores

- Semaphore with an **initial count of 1**
- Can be used as a **mutex** with non-reentrant locking semantics
  - Whoever holds the sole permit holds the mutex

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.8

## Using Semaphores to bound a collection

```
public BoundedHashSet<T> {  
    private final Set<T> set;  
    private final Semaphore sem;  
    public BoundedHashSet(int bound) {  
        this.set = Collections.synchronizedSet(new HashSet<T>());  
        sem = new Semaphore(bound);  
    }  
    public boolean add(T o) throws InterruptedException {  
        sem.acquire();  
        boolean wasAdded = false;  
        try {  
            wasAdded = set.add(o);  
            return wasAdded;  
        } finally {  
            if (!wasAdded) sem.release();  
        }  
    }  
    public boolean remove(Object o) {  
        boolean wasRemoved = set.remove(o);  
        if (wasRemoved) sem.release();  
        return wasRemoved;  
    }  
}
```

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

L11.9

## Barriers

- Barriers are similar to latches in that they **block a group of threads** till an event has occurred
- All threads must come together at **barrier point at the same time** to proceed
  - Latches wait for events, barriers **wait for other threads**

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.10

## Barriers and dinner ...

- Family rendezvous protocol
- Everyone meet at Panera @ 6:00 pm;
  - Once you get there, stay there ... till everyone shows up
  - Then we'll figure out what we do next

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.11

## Barriers

- Often used in simulations where work to calculate one step can be done in parallel
  - But all work associated with a given step must complete before advancing to the next step
- All threads complete step  $k$ , before moving on to step  $k+1$

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.12

## CyclicBarrier

- Allows a fixed number of parties to *rendezvous* at a fixed point
- Useful in **parallel iterative algorithms**
  - Break problem into fixed number of independent subproblems
- Creation of a CyclicBarrier
  - Runnable cyclicBarrierAction = ... ;  
 cyclicBarrier cyclicBarrier =  
     new CyclicBarrier(2, cyclicBarrierAction);

February 20, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L11.13

## Using Cyclic Barriers

Source: From the Java API

```
class Solver {
    final int N;    final CyclicBarrier barrier;
    class Worker implements Runnable {
        int myRow;
        Worker(int row) { myRow = row; }
        public void run() {
            while (!done()) {
                processRow(myRow);
                try {
                    barrier.await();
                } catch (BrokenBarrierException ex) {
                    ...
                }
            }
        }
    }
    public Solver(float[][] matrix) {
        data = matrix;    N = matrix.length;
        barrier = new CyclicBarrier(N, new Runnable() { public void run() {
            mergeRows(...); } });
        for (int i = 0; i < N; ++i)
            new Thread(new Worker(i)).start(); //DO NOT START THREAD in constructor.
        waitUntilDone();
    }
}
```

February 20, 2018  
 Instructor: SHRIDEEP PALICKARA

L11.14

## Exchanger

- Another type of barrier
- Two-party barrier
- Parties **exchange data** at the barrier point
- Useful when asymmetric activities are performed
  - Producer-consumer problem
- When 2 threads exchange objects via Exchanger
  - Safe publication of objects to other party

February 20, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L11.15

## THREAD SAFETY SUMMARY

February 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L11.16

## Thread Safety: Summary

[1/4]

- It's all about *mutable, shared state*
  - The less mutable state there is, the easier it is to ensure thread-safety
- Make fields **final** unless they need to be mutable
- **Immutable** objects are automatically thread-safe
- **Encapsulation** makes it practical to manage complexity

February 20, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L11.17

## Thread Safety: Summary

[2/4]

- Guard each mutable variable with a **lock**
- Guard all variables in an invariant with the **same lock**
- Hold locks for the *duration* of compound actions

February 20, 2018  
 Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
 Dept. Of Computer Science, Colorado State University

L11.18

## Thread Safety: Summary

[3/4]

- Program that access mutable variables from multiple threads without synchronization?
  - Broken program
- Include thread-safety in the design process
  - Document if your class is not thread-safe
- Document your synchronization policy

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.19

## Thread Safety: Summary

[4/4]

- Rather than scattering access to shared state throughout your programs and attempting *ad hoc* reasoning about interleaved access
  - Structure program to facilitate reasoning about concurrency
  - Use a set of standard synchronization primitives to control access to shared state

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.20


## MAPREDUCE

February 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.21

## MapReduce: What we will look at



February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.22

## CLOUD COMPUTING

February 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.23

## The volume of data that we produce has increased dramatically

- IDC (International Data Corporation) estimates
  - 180 EB ( $10^{18}$ ) in 2006
  - 1.8 ZB ( $10^{21}$ ) in 2011
    - Roughly a disk drive per person!
  - 40 ZB by 2020

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.24

### Some of the sources of this deluge

- New York Stock Exchange
  - 1 TB of new trade data every day
- Facebook
  - $\sim 10^{12}$  photos
- Internet Archive
  - Stores 2 PB of data ... growing at 20 TB per month
- LHC produces 15 PB per year

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.25

### Amount of data generated by machines will outpace what people produce

- Machine logs
- RFID readers
- Sensor networks
- Instruments
- Vehicle GPS traces
- IoT
  - 20-35 billion IoT devices are expected to be online in 2020

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.26

### Hard disk capacities, seek rates, and transfer times

- 1990
  - 1 GB HDDs with a transfer speed of 4.4 MB/sec
- Now
  - 1 TB hard drives are common
  - But the transfer speed is just 100 MB/sec
  - Writing is even slower!

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.27

### Data transfers can be improved by using multiple disks

- What if we use 100 disk drives?
  - Each holding 1/100<sup>th</sup> of the data
- We could have **cumulative transfer** speeds of up to 100 x 100 MB/sec or 10 GB/sec
- But isn't using 1/100<sup>th</sup> of disk wasteful?
  - Not if you store a 100 different datasets on these disks
  - Provide shared access to the disks

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.28

### But there's more than just reading and writing from multiple disks in parallel

- **Cope with hardware failures**
  - As the number of components increase, so does the probability of failure
- Analysis tasks need to be able to **combine data**
  - Dataset is dispersed over multiple disks

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.29

### What MapReduce provides ...

- Programming model that **abstracts** the problem from disk reads and writes
- Transform the problem into **computations** over sets of keys and values
- Supports **distributed processing** on large datasets over a cluster of computers

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.30

### But why not use databases with lots of disks? [1/2]

- Another trend in disk drives
  - ▣ Seek time is improving *much slower* than transfer rates
- If data access pattern is dominated by seeks?
  - ▣ It takes longer to read or write large portions of the dataset than streaming through it
    - Streaming through dataset operates at transfer speed

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.31

### But why not use databases with lots of disks? [2/2]

- Updating a small proportion of records in the dataset
  - ▣ Traditional B-Tree works well
- For updating a majority of the dataset
  - ▣ B-Tree is less efficient than MapReduce which uses Sort/Merge to rebuild the dataset

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.32

### MapReduce should be seen as being complementary to databases

- MapReduce is good for problems that access the **entire dataset**
  - ▣ Particularly *ad hoc* analysis
  - ▣ Write once, read many times
- RDBMS is good for point queries or updates
  - ▣ Dataset **has been indexed** for low-latency retrieval and update times
  - ▣ Read and write many times

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.33

### Grid Computing/HPC systems

- Distribute work across a cluster of machines that access a **shared file system**
- Works well for predominantly compute-intensive jobs
  - ▣ Problem when access to large data volumes is needed
    - Network bandwidth is a bottleneck and compute nodes become idle

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.34

### MapReduce tries to colocate data with the compute node

- **Data Locality**
  - ▣ Data access is fast since it is local
  - ▣ Conserves network bandwidth
- Implementations go to great lengths to conserve it
  - ▣ Model network topology

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.35

### MPI (Message Passing Interface) gives great control to the programmer

- MPI requires **explicit handling** of the mechanics of *data flow*
  - ▣ In MapReduce, the mechanics of data flow is implicit
- MapReduce spares programmers from having to think about failures
  - ▣ Detect failures and schedule replacements on healthy machines
  - ▣ Done with a **shared-nothing architecture**
  - ▣ MPI programs have to deal with checkpointing and recovery
    - More control but difficult to write

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.36

## Volunteer computing

- SETI@home
- Volunteers donate cycles not bandwidth
- MapReduce
  - ▣ Runs jobs lasting minutes or hours on trusted, dedicated machines with **high-bandwidth** interconnects
- Volunteer computing
  - ▣ Perpetual computations on untrusted machines
    - **Highly variable connection** speeds and no data locality

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.37

## MAPREDUCE

MATERIALS BASED ON  
JEFFREY DEAN and SANJAY GHEMAWAT: *MapReduce: Simplified Data Processing on Large Clusters*. OSDI 2004: 137-150

February 20, 2018

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.38

## Source of raw data at Google

- **Crawled** data
- **Log** of the web requests

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.39

## Several computations work on this raw data to compute derived data

- Inverted indices
- Representation of the graph structure of web documents
- Pages crawled per host
- Most frequent queries in a day ...

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.40

## Most computations are conceptually straightforward

- But data is large
- Computations must be **scalable**
  - ▣ Distributed across thousands of machines
  - ▣ To complete in a reasonable amount of time

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.41

## Complexity of managing distributed computations can ...

- Obscure **simplicity** of original computation
- Contributing factors:
  - ▣ How to **parallelize** the computation
  - ▣ Distribute the **data**
  - ▣ Handle **failures**

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.42

## MapReduce was developed to cope with this complexity

- Express simple computations
- Hide messy details of:
  - ① Parallelization
  - ② Data distribution
  - ③ Fault tolerance
  - ④ Load balancing

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.43

## MapReduce

- Programming model
- Associated implementation for
  - Processing & Generating large data sets

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.44

## Programming model

- Computation takes a set of **input** key/value pairs
- Produces a set of **output** key/value pairs
- Express the computation as two functions:
  - Map
  - Reduce

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.45

## Map

- Takes an input pair
- Produces a set of intermediate key/value pairs

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.46

## Mappers

- If map operations are **independent** of each other they can be performed in parallel
  - **Shared nothing**
- This is usually the case

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.47

## MapReduce library

- **Groups** all intermediate values with the same intermediate key
- **Passes** them to the Reduce function

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.48



## Reduce function

- Accepts intermediate key **k** and
  - Set of values for that key
- **Merge** these values together to get
  - Smaller set of value

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.49

## Counting number occurrences of each word in a large collection of documents

```
map (String key, String value)
    //key: document name
    //value: document contents

    for each word w in value
        EmitIntermediate(w, "1")
```

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.50

## Counting number occurrences of each word in a large collection of documents

```
reduce (String key, Iterator values)
    //key: a word
    //value: a list of counts

    int result = 0;
    for each v in values
        result += ParseInt(v);
    Emit(AsString(result));
```

Sums together all counts  
emitted for a particular word

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.51

## The contents of this slide set are based on the following references

- Hadoop: The Definitive Guide by Tom White. Early Release. 3<sup>rd</sup> Edition. O'Reilly. [Chapter 1]
- Jeffrey Dean, Sanjay Chawala: MapReduce: Simplified Data Processing on Large Clusters. OSDI 2004: 137-150
- Jeffrey Dean, Sanjay Chawala: MapReduce: simplified data processing on large clusters. Commun. ACM 51(1): 107-113 (2008)

February 20, 2018  
Instructor: SHRIDEEP PALICKARA

CS455: Introduction to Distributed Systems [Spring 2018]  
Dept. Of Computer Science, Colorado State University

L11.52